


TechGI IV - ComDis


ComDis

Introduction to
Communication Networks
and **D**istributed Systems



***Unit 2: Reference Models
and beyond...***

Internet maps



Large Networks need structure! Why?

- Scaling

- Remember: each switch knows route to each destination...
- Hierarchy usually simplifies a lot...

- Locality


- Close hosts are clustered,
Local networks

- Heterogeneity


- Different applications (e.g. control, sensing) have requirements best served by NNN
- Multiple technologies used....

- Administration

- Who sets the rules for usage ???



Internet: Interoperability vs. Heterogeneity



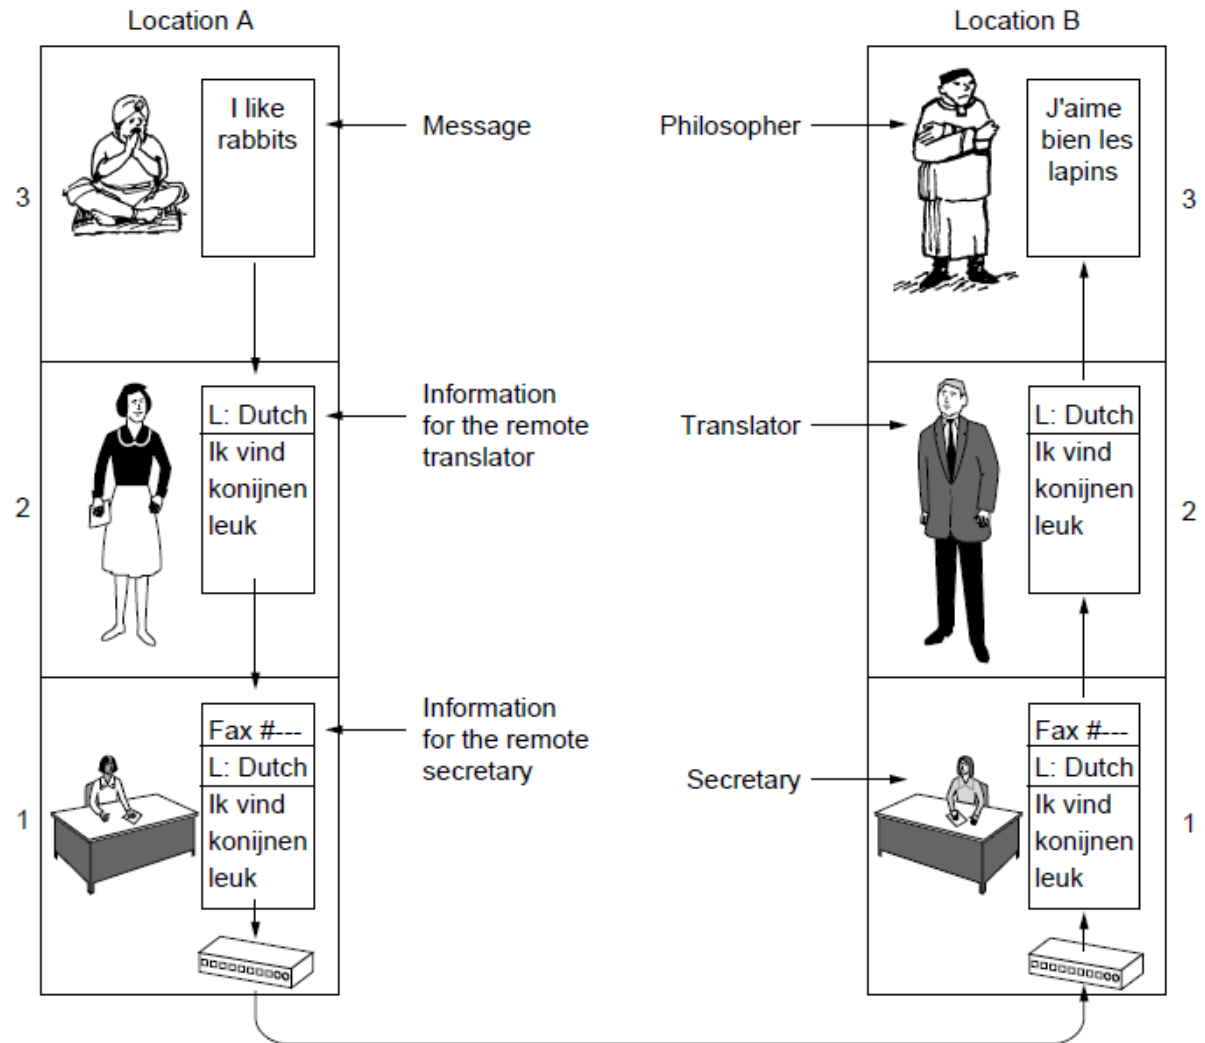
See: [The Brief History of the Internet](#)

Do you remember the philosophers...

Philosophers

Assistants

Office clerks



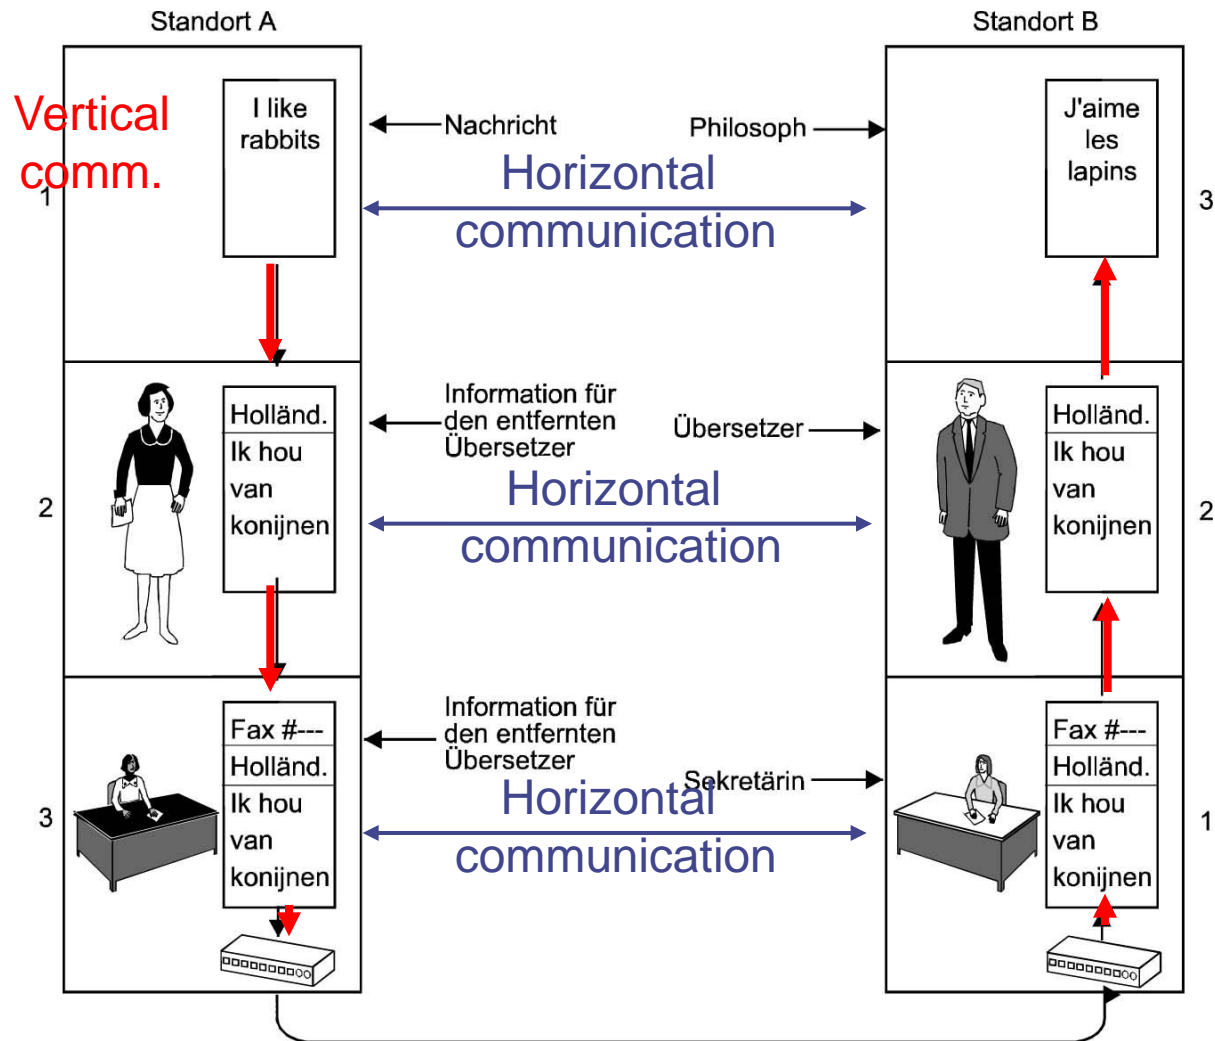
The reference model

- To keep complexity of communication systems tractable, division in subsystems with clearly assigned responsibilities is necessary - layering
- Each layer (and the communication system as a whole) offers a particular service
 - Services become more abstract and more powerful the higher up in the layering hierarchy
- To provide a service, a layer has to be distributed over remote devices
- Remote parts of a layer use a protocol to cooperate
 - Make use of service of the underlying layer to exchange data
 - Protocol is a horizontal relationship, service a vertical relationship
- Layers/protocols are arranged as a (protocol) stack
 - One atop the other, only using services from directly beneath

⇒ Strict layering

Analogy: Nested layers as nested translations

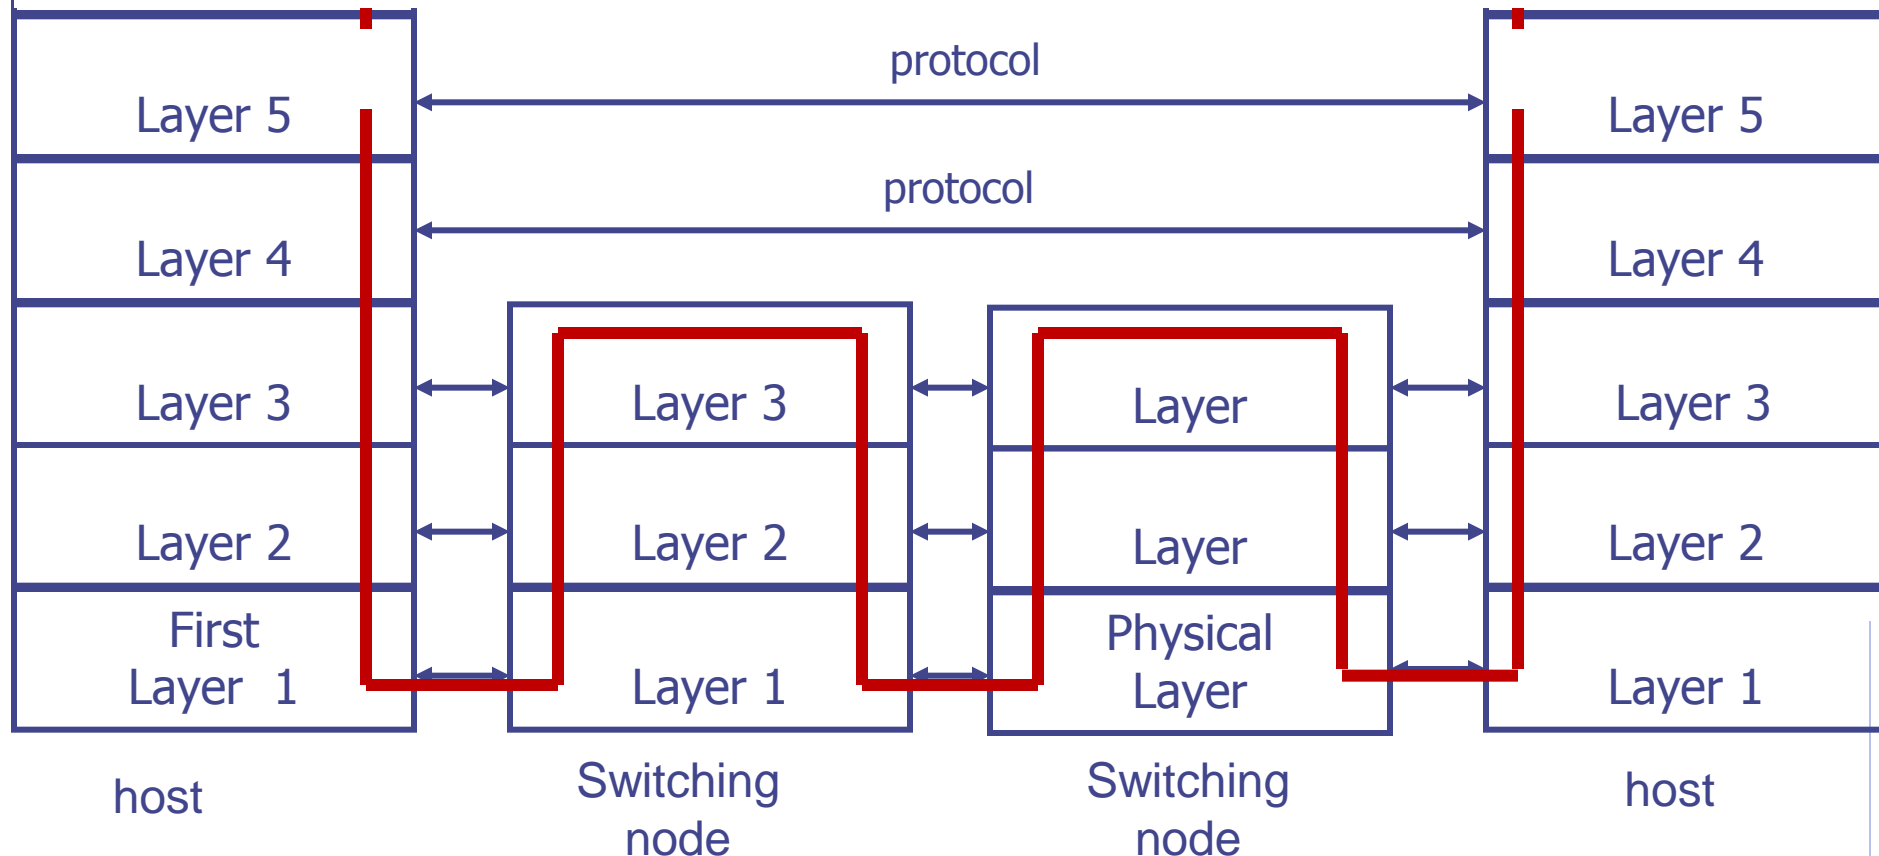
- Vertical vs. horizontal communication
 - Vertical: always real
 - Horizontal: may be real or virtual



Horizontal (real!) communication

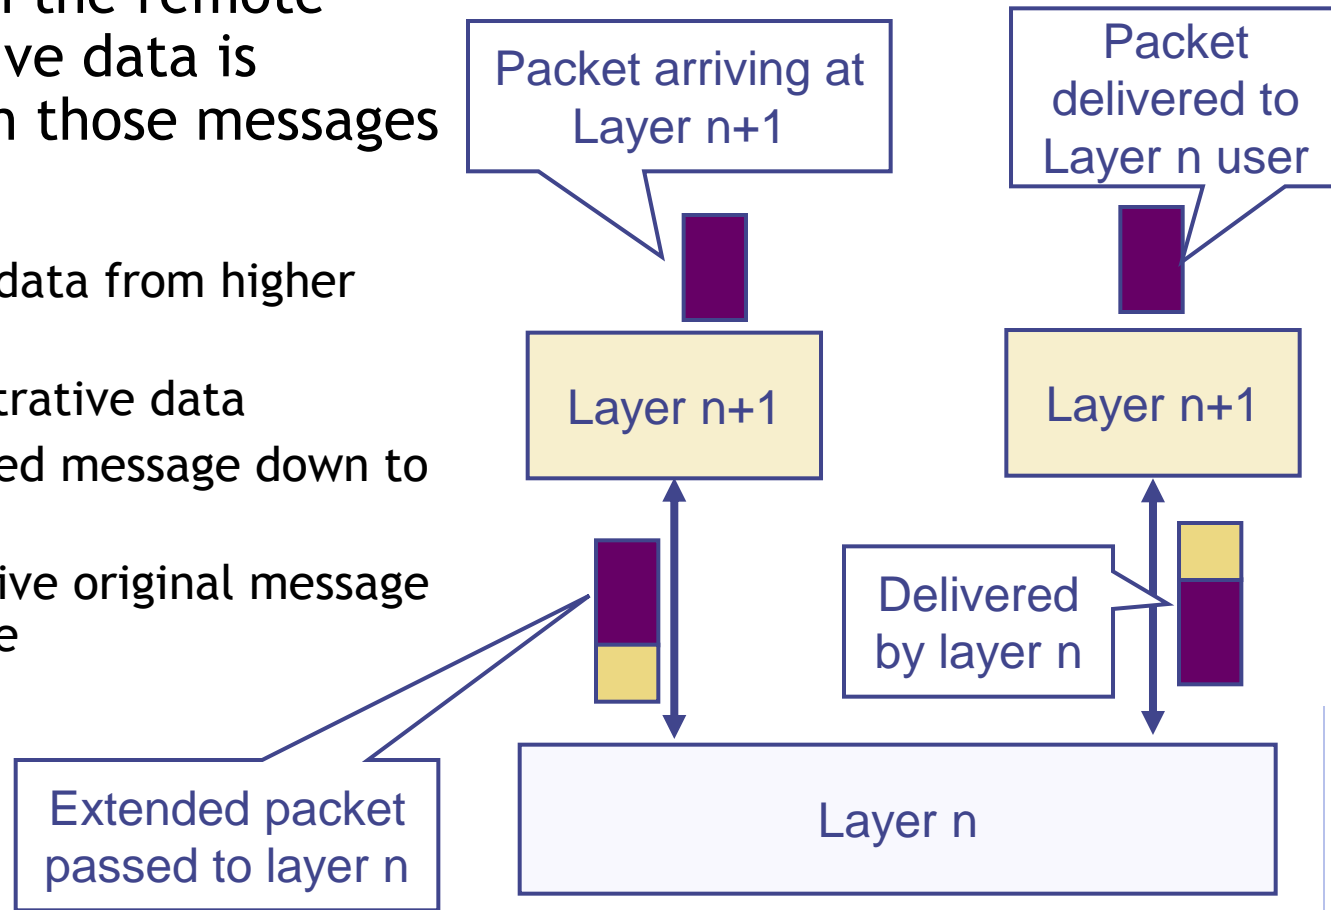
Multi-layer Architecture

- Number of Layers, and { services, naming and addressing conventions) / Layer
- Functions to be executed in each layer
- Protocols: (host- to – host, node –to – node, Host-to-node)



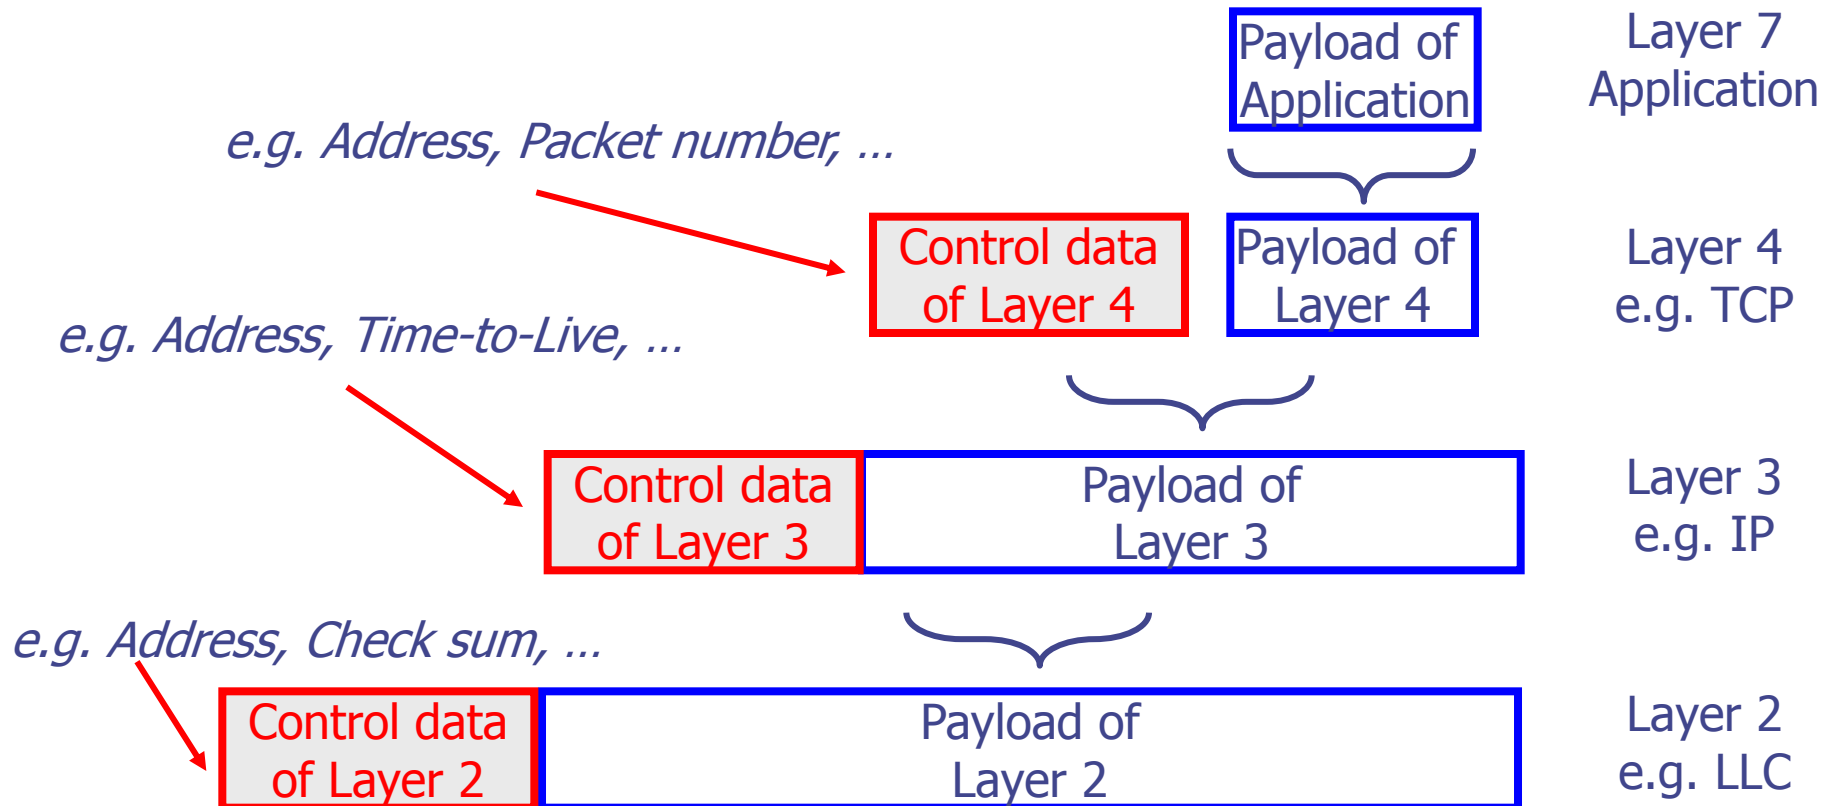
Protocols and messages

- When using lower-layer services to communicate with the remote peer, administrative data is usually included in those messages
- Typical example
 - Protocol receives data from higher layer
 - Adds own administrative data
 - Passes the extended message down to the lower layer
 - Receiver will receive original message plus administrative data
- Encapsulating
 - Header or trailer



Embedding messages

- Messages from upper layers are used as payload for messages in lower layers



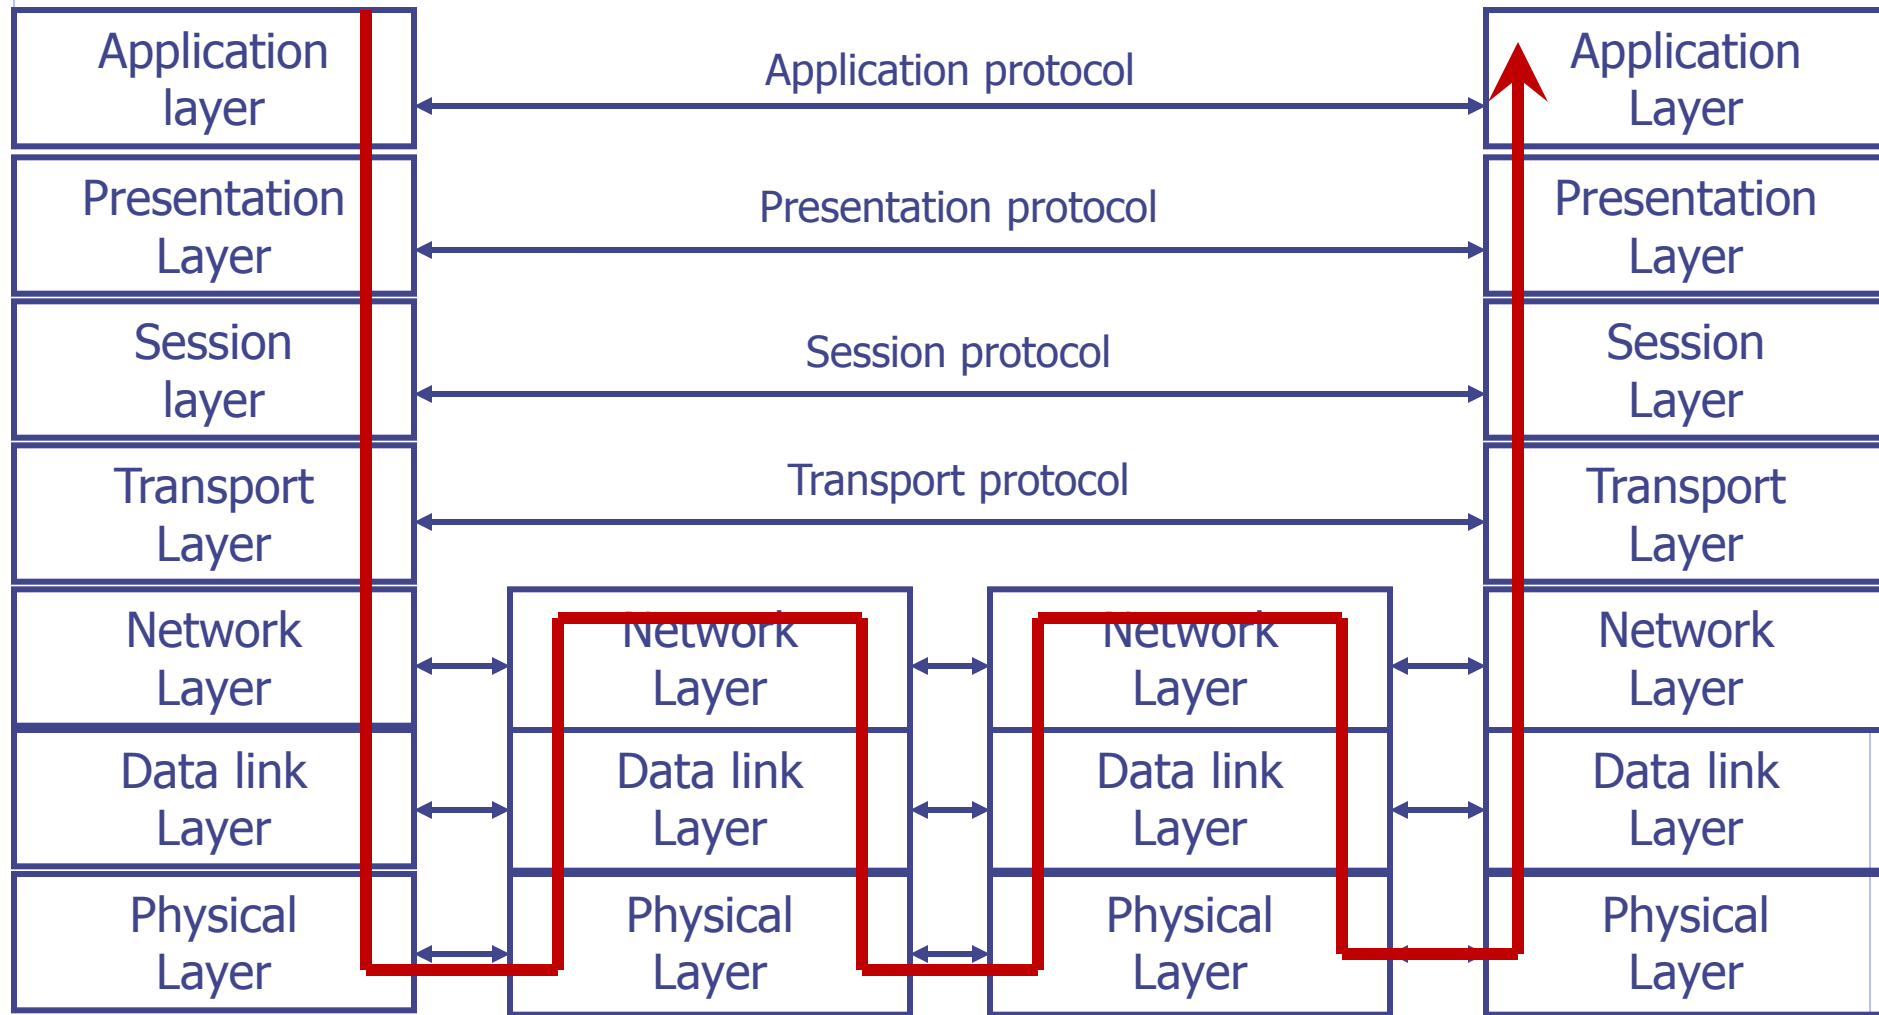
How to structure functions/layers ?

- Many functions have to be realized
- Not each function is necessary in each Layer..
- How to actually assign them into layers to obtain a real, working communication system?
 - This is the role of a specific reference model
- Two main reference models exist
 - ISO/OSI reference model (International Standards Organization Open Systems Interconnection)
 - TCP/IP reference model (by IETF - Internet Engineering Taskforce)

ISO/OSI reference model

- Basic design principles
 - One layer per abstraction of the “set of duties”
 - Choose layer boundaries such that information flow across the boundary is minimized (minimize inter-layer interaction)
 - Enough layers to keep separate things separate, few enough to keep architecture manageable
- Result: 7-layer model
 - Not strictly speaking an architecture, because
 - Precise interfaces are not specified (nor protocol details !)
 - Only general duties of each layer are defined

ISO/OSI model



7 layers in brief

- Physical layer: Transmit raw bits over a physical medium
- Data Link layer: Provide a (more or less) error-free transmission service for data frames - also over a shared medium!
- Network layer: Solve the forwarding and routing problem for a network- bring data to a **desired host**
- **Transport layer**: Provide (possibly reliable, in order) end-to-end communication, overload protection, fragmentation to **processes**

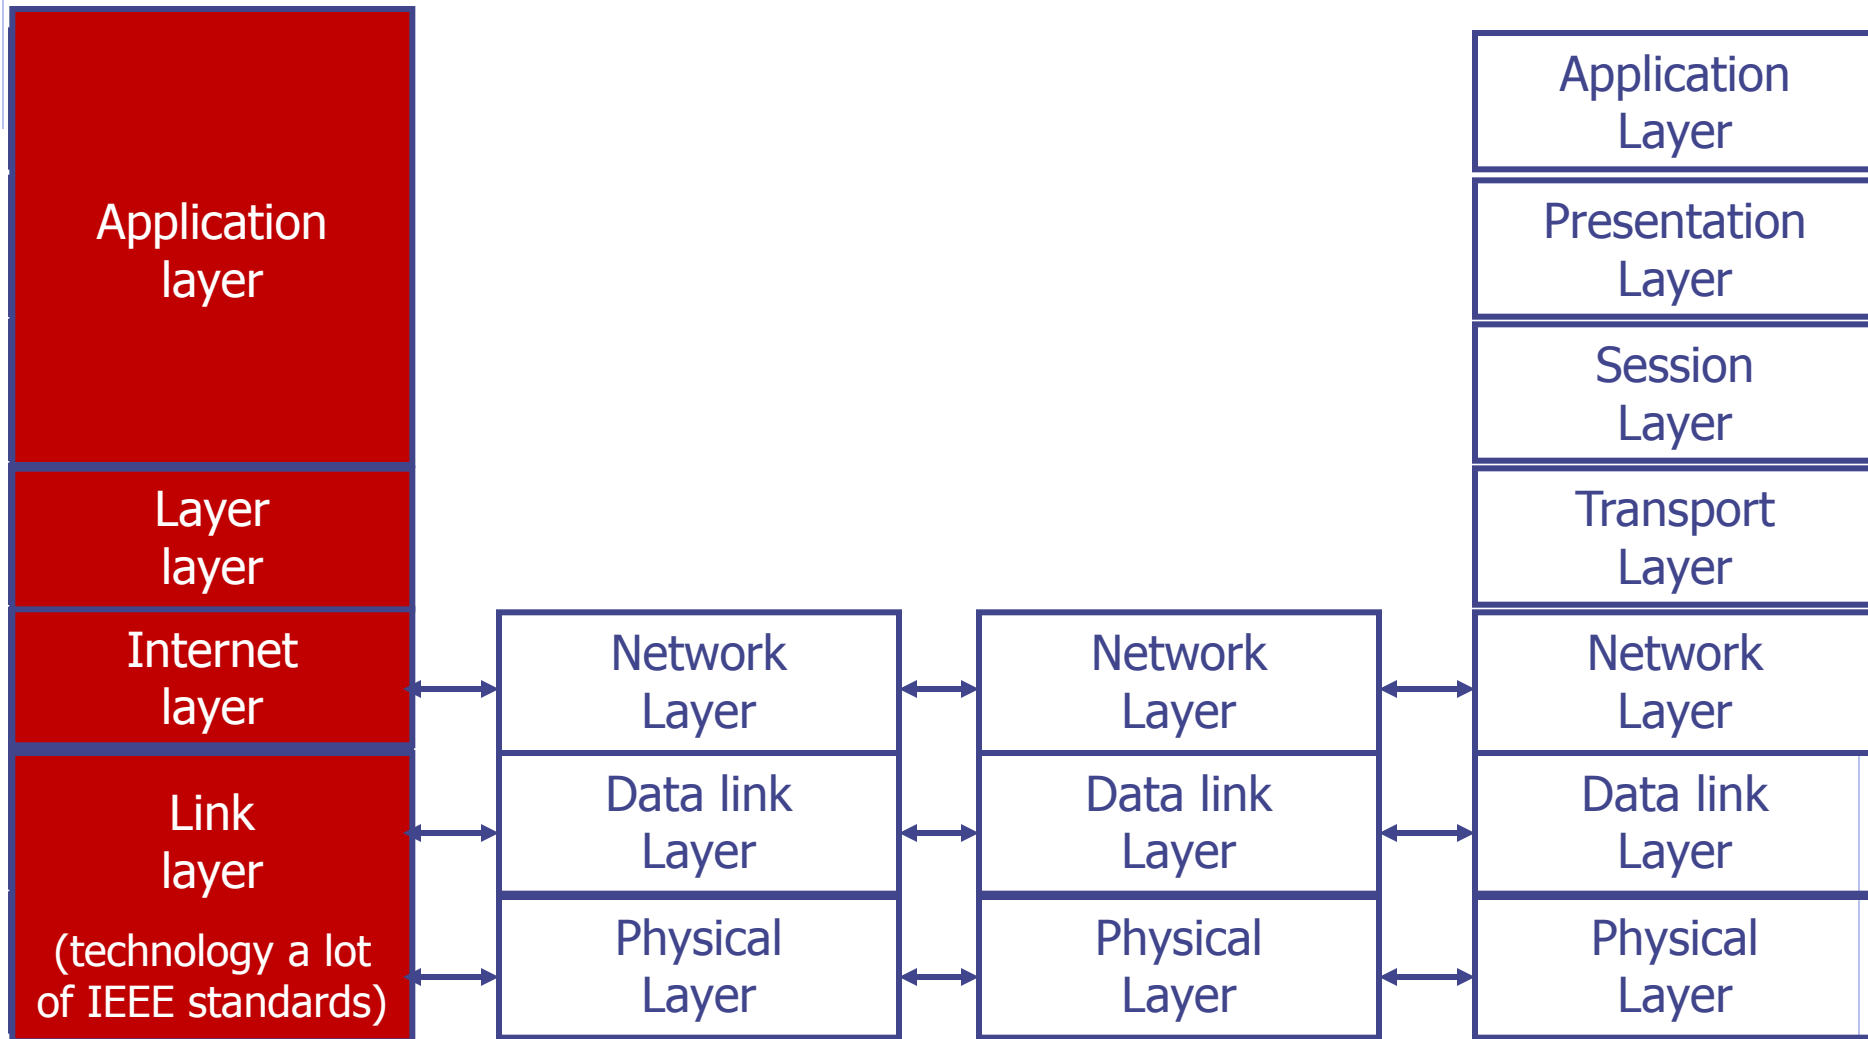
“Bringing data from process A to B with sufficient quality”

- Session layer: Group communication into sessions which can be synchronized, checkpointed, ...
- Presentation layer: Ensure that syntax and semantic of data is uniform between all types of terminals
- Application layer: Actual application, e.g., protocols to transport Web pages

TCP/IP reference model

- Historically based on ARPANET, evolving to the Internet
 - Started out as little university networks, which had to be interconnected
- Some generic rules & principles
 - Internet connects **networks**
 - Minimum functionality assumed (just unreliable packet delivery!)
 - Internet layer (IP): packet switching, addressing, routing & forwarding
 - ➔ *Internet over everything*
 - End-to-end
 - Any functionality should be pushed to the instance needing it!
 - Fate sharing
- In effect only two layers really defined... Internet and Transport Layer - Lower and higher layers not really defined
 - ➔ *Anything over internet*
- New Applications do NOT need any changes in the NETWORK!
 - Compare with the telephone network!!!

TCP/IP model (in red)



Some example protocols

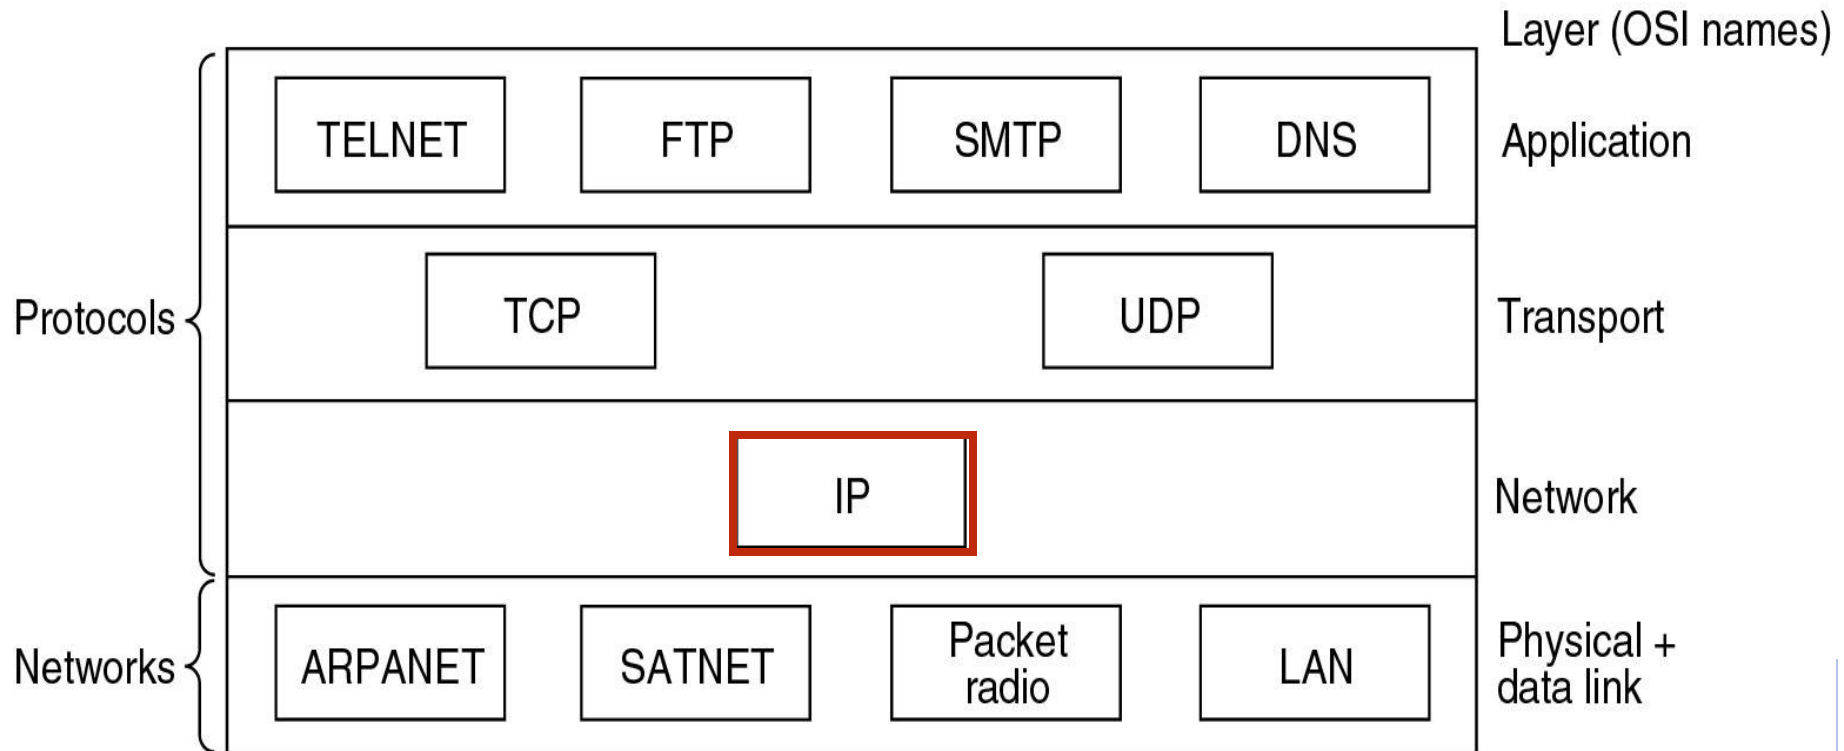
- A communication architectures needs standard protocols in addition to a layering structure
- And: some generic rules & principles which are not really a protocol but needed nonetheless
 - Example principle: end-to-end
 - Example rule: Naming & addressing scheme
- Popular protocols of the 5-layer reference model
 - Data link layer: Ethernet & CSMA/CD
 - Network layer: Internet Protocol (IP)
 - Transport layer: Transmission Control Protocol (TCP)

Standardization

- To build large networks, standardization is necessary
- Traditional organization from a telecommunication/ telephony background
 - Well established, world-wide, relatively slow “time to market”
- Internet
 - Mostly centered around the Internet Engineering Task Force (IETF) with associated bodies (Internet Architectural Board, Internet Research Task Force, Internet Engineering Steering Group)
 - Consensus oriented, focus on working implementations
 - Hope is quick time to market, but has slowed down considerably in recent years
- Manufacturer bodies

TCP/IP – Suite of protocols

- Over time, suite of protocols evolved around core TCP/IP protocols



So-called “**hourglass model**”: Thin waist of the protocol stack at IP, above the technological layers


Naming & addressing in the TPC/IP stack

- Names: Data to identify an entity exist on different levels
 - Alphanumerical names for machines: www.tu-berlin.de
- Address: Data how/where to find an entity
 - Address of a network device in an IP network: An IP address
 - IPv4: 32 bits, structured into 4x8 bits
 - Example: 131.234.20.99 (dotted decimal notation)
 - Address of a network: Some of the initial bits of an IP address
- Address of a networked device in the LOCAL AREA (IEEE 802 standardized) network...
 - 48 bits, hexadecimal notation, example: 08:00:20:ae:fd:7e

Mapping

- Needed: Mapping from name to address
 - ➔ Realized by separate protocols
- From alphanumerical name to IP address: Domain Name System (DNS)
- Often also useful: Mapping from IP address to MAC name/address: Address Resolution Protocol (ARP)


Web server process'
service access point



Understanding Ports

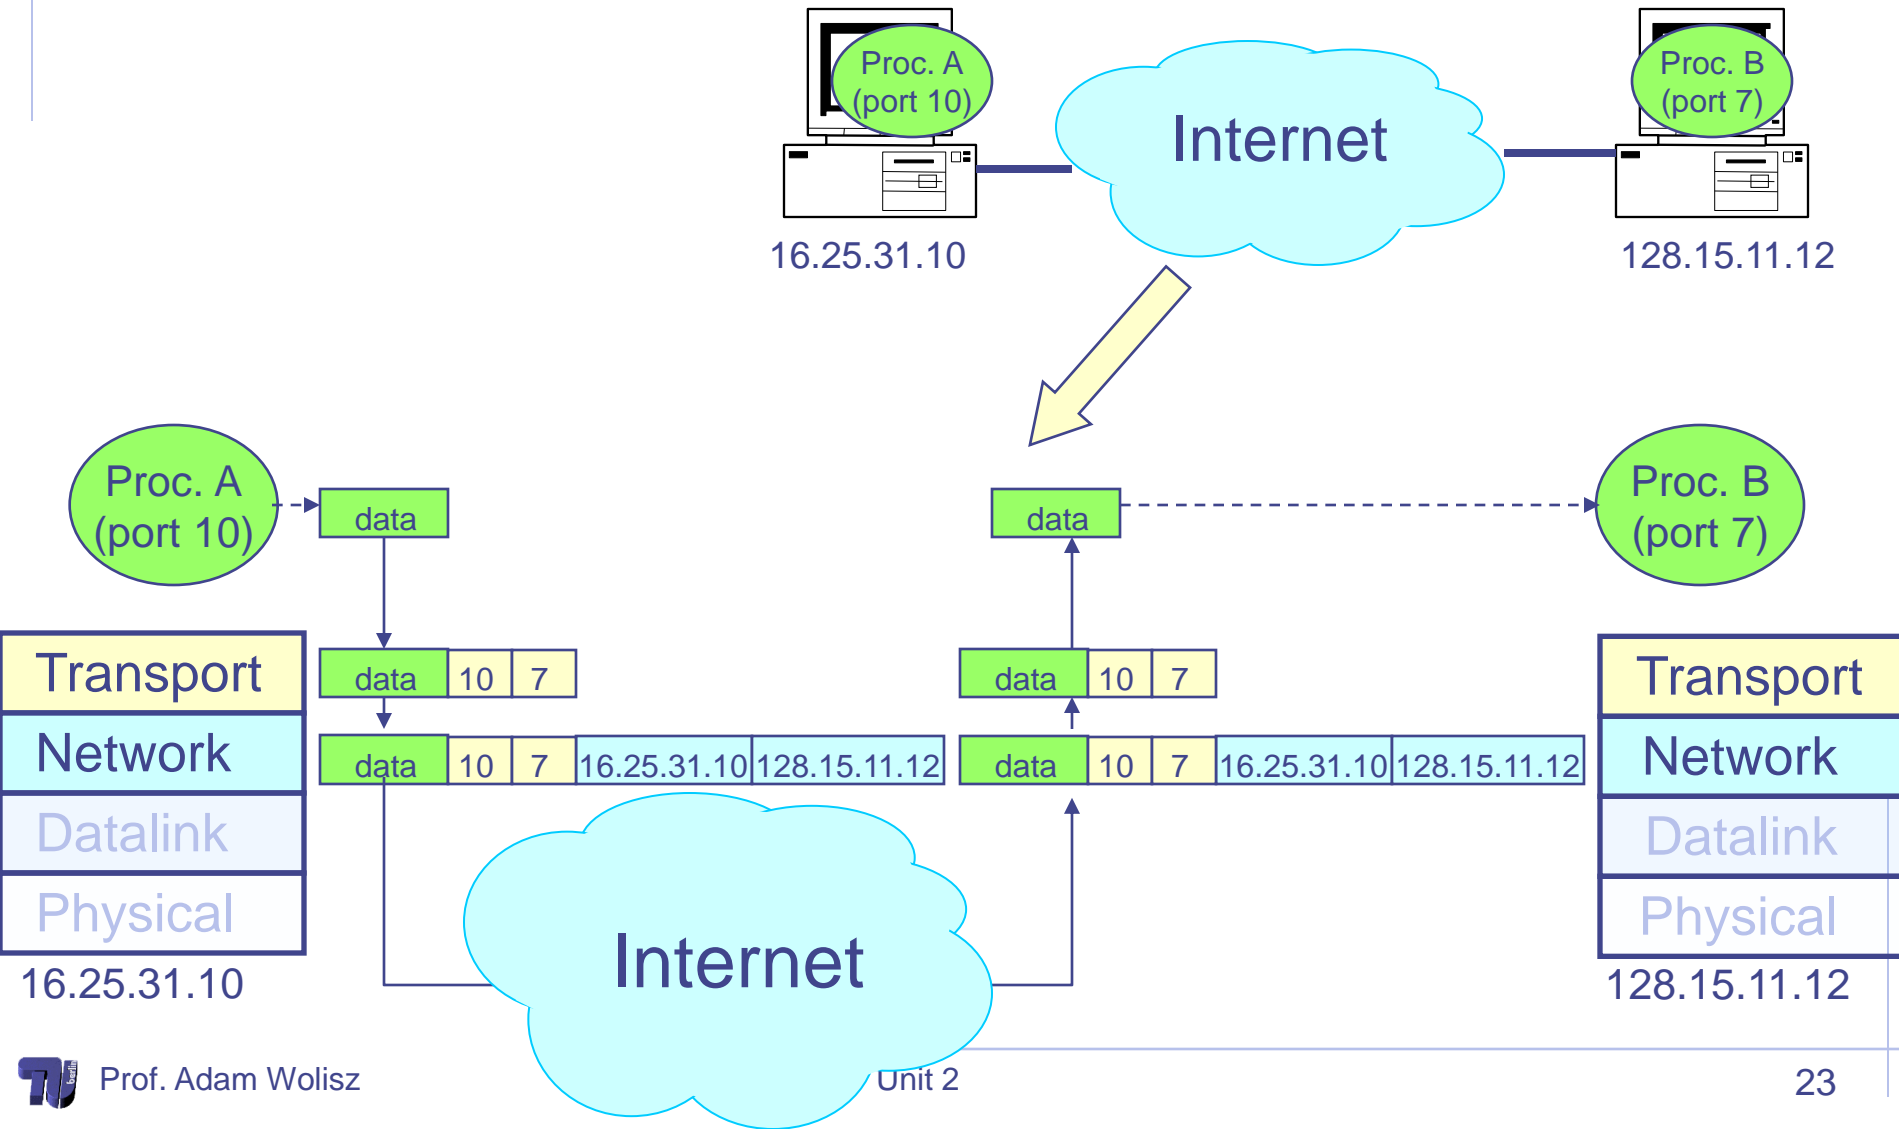
[Buya]

- Port is represented by a positive (16-bit) integer value
- Some ports have been reserved to support common/well known services: http 80/tcp; ftp 21/tcp; telnet 23/tcp; smtp 25/tcp;
- User level process/services generally use port number value ≥ 1024




End-to-End Layering View

Stoica

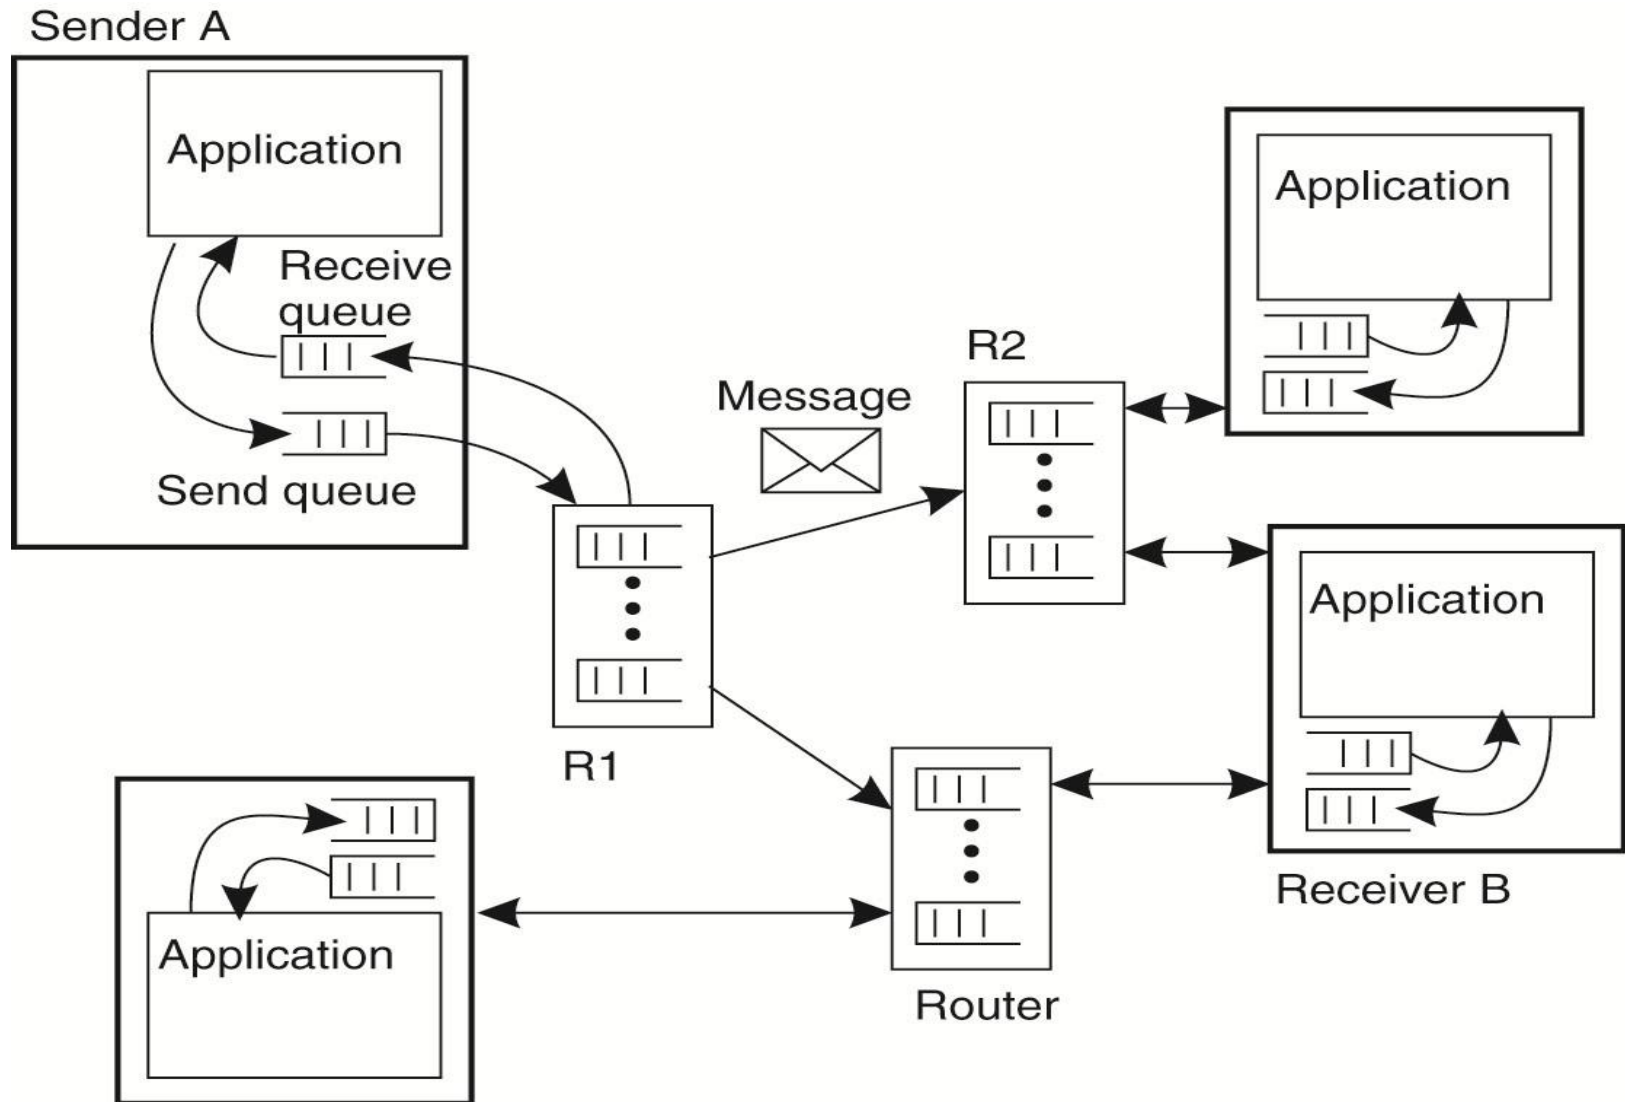


- Process A sends a packet to process B



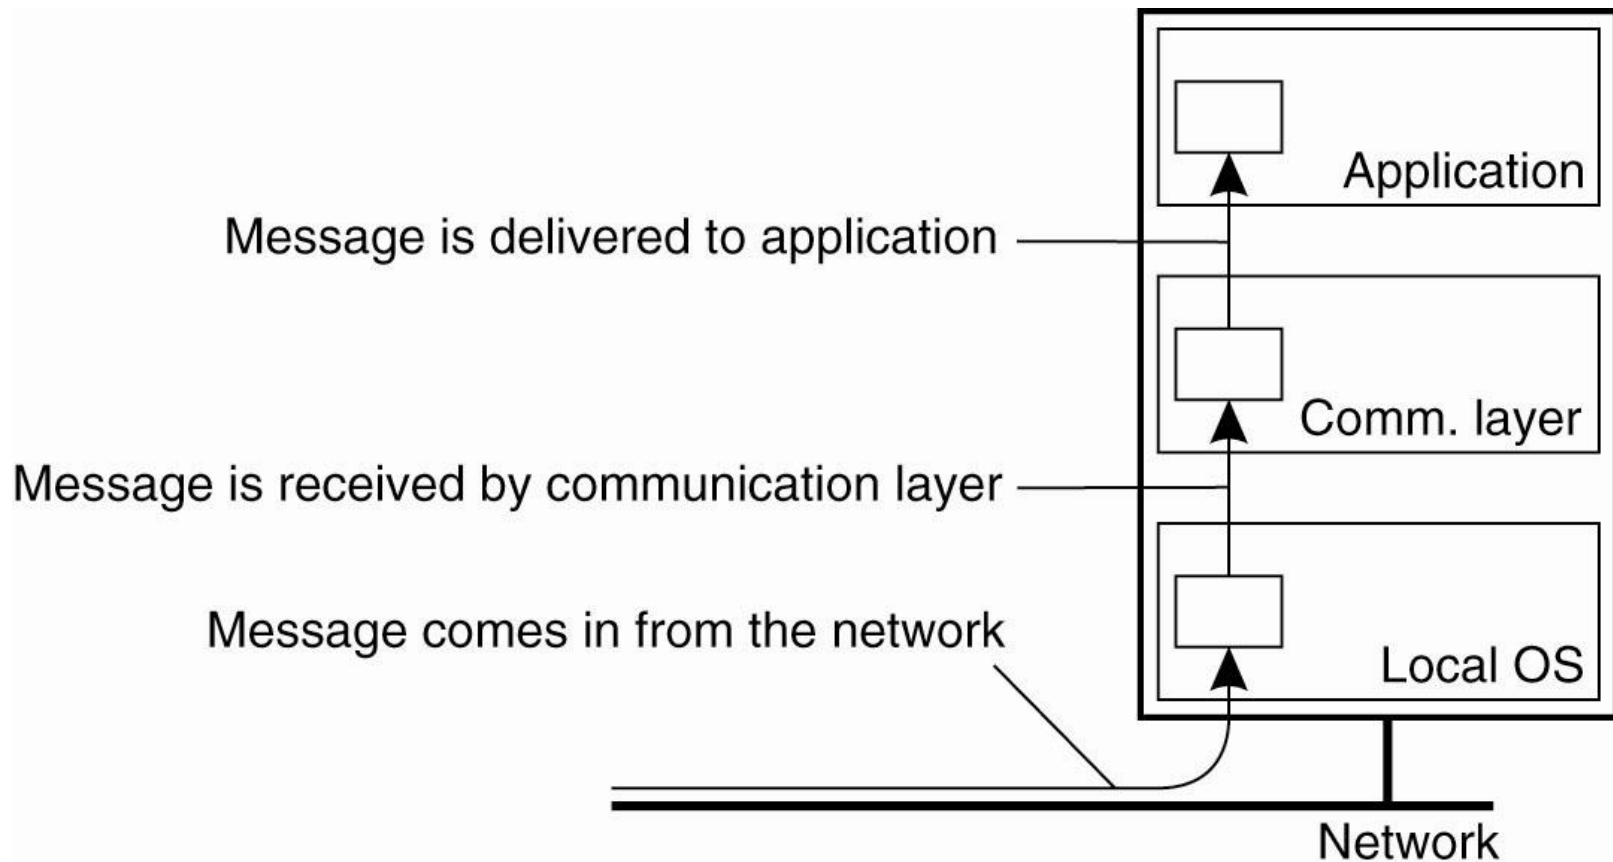
IP Address:

A four-part “number” used by *Network Layer* to route a packet from one computer to another



Message Receipt vs. Message Delivery [Tanenbaum]

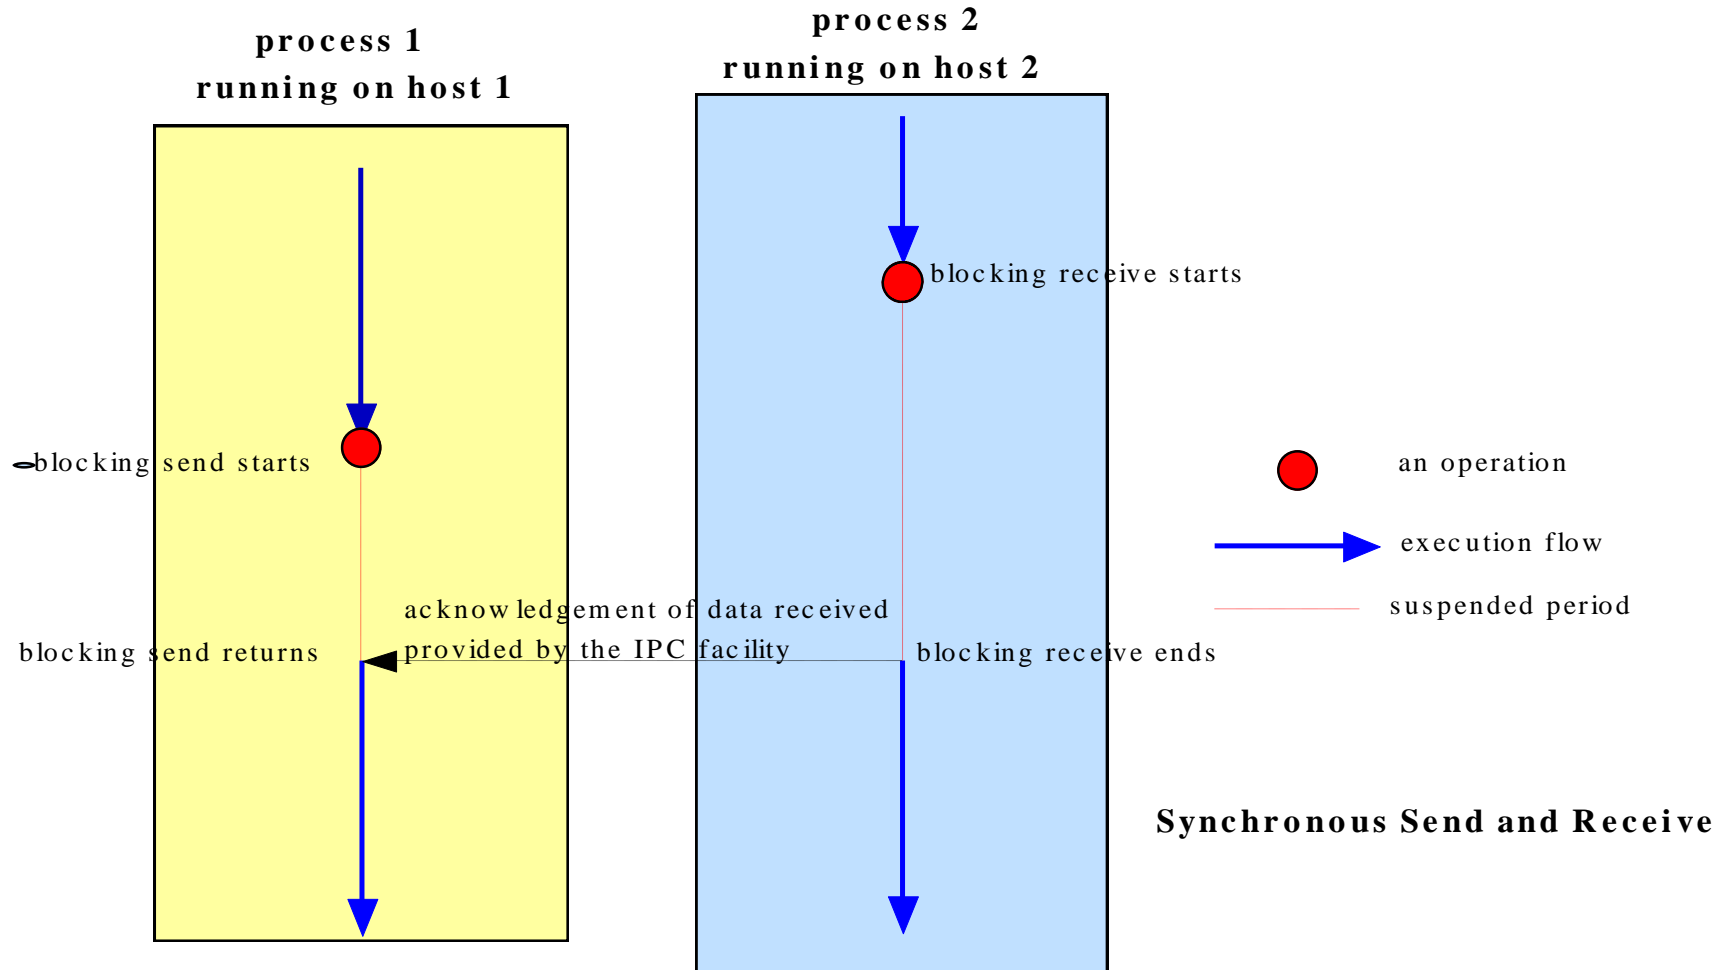
- Figure 8-12. The logical organization of a distributed system to distinguish between message receipt and message delivery.



Synchronous Interaction

- Blocking send
 - Blocks until message is transmitted
 - Blocks until message acknowledged
- Blocking receive
 - Waits for message to be received
- Known upper/lower bounds on execution speeds, message transmission delays and clock drift rates
- More difficult to build, conceptually simpler model
 - use Queue (for waiting)
 - *send* and *receive* are blocking

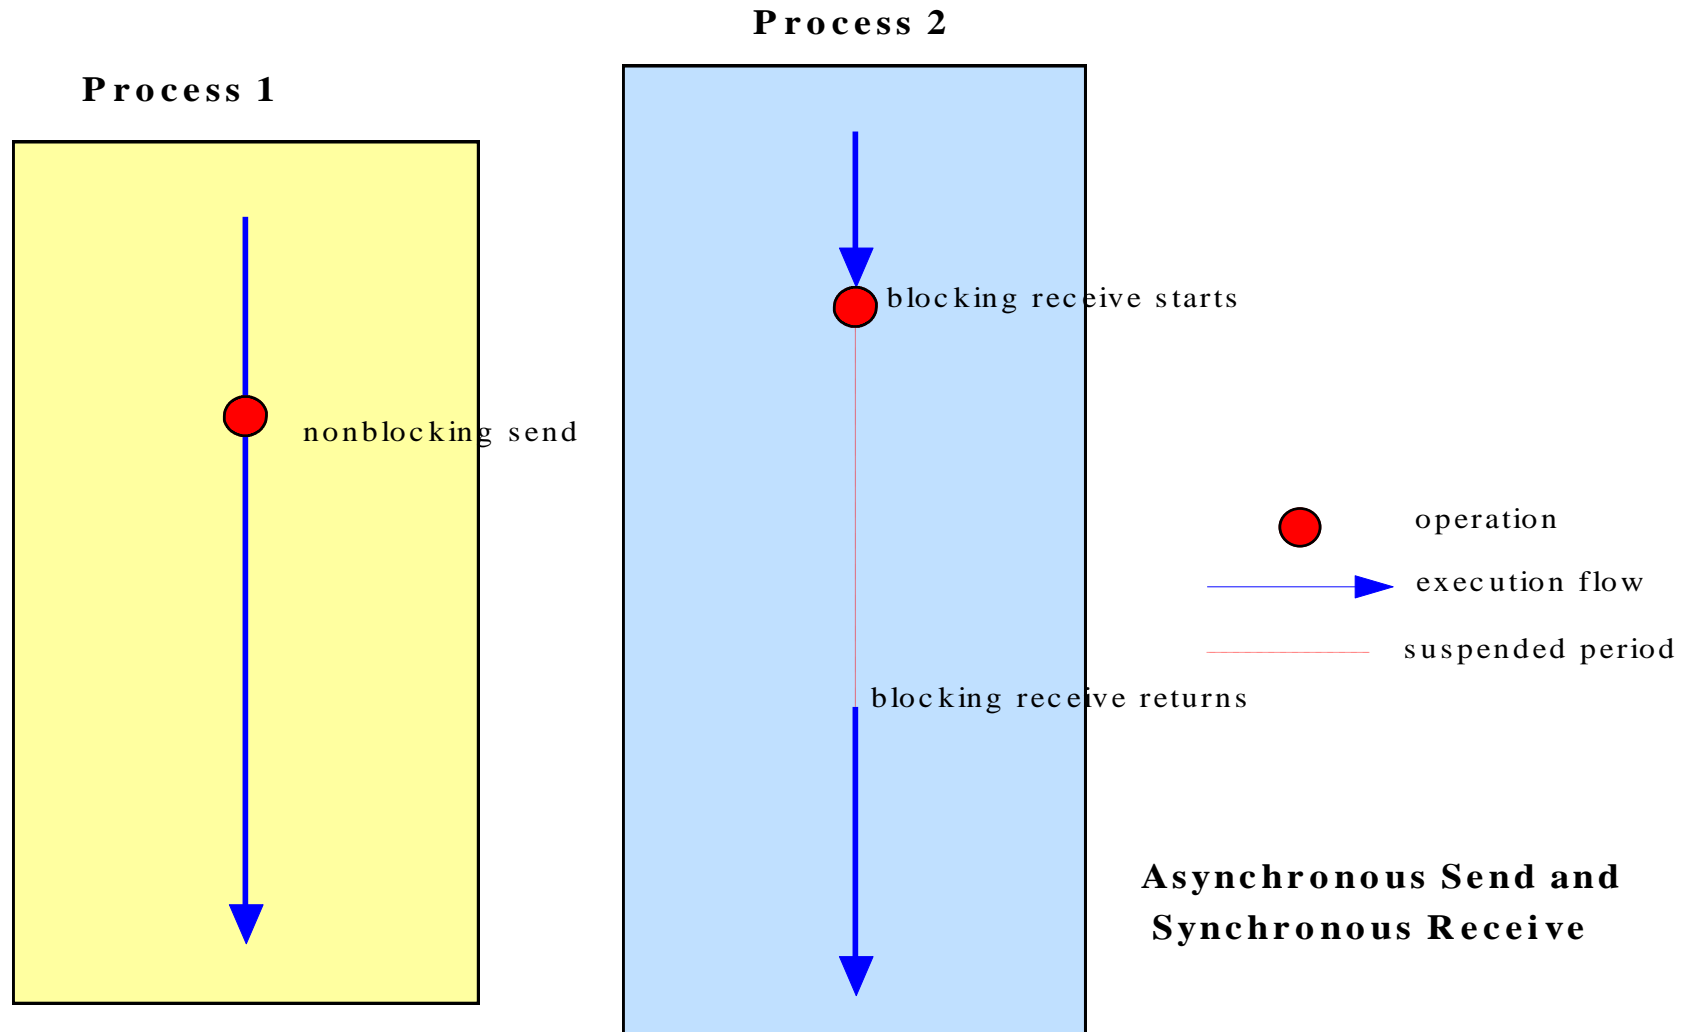
Synchronous send and receive



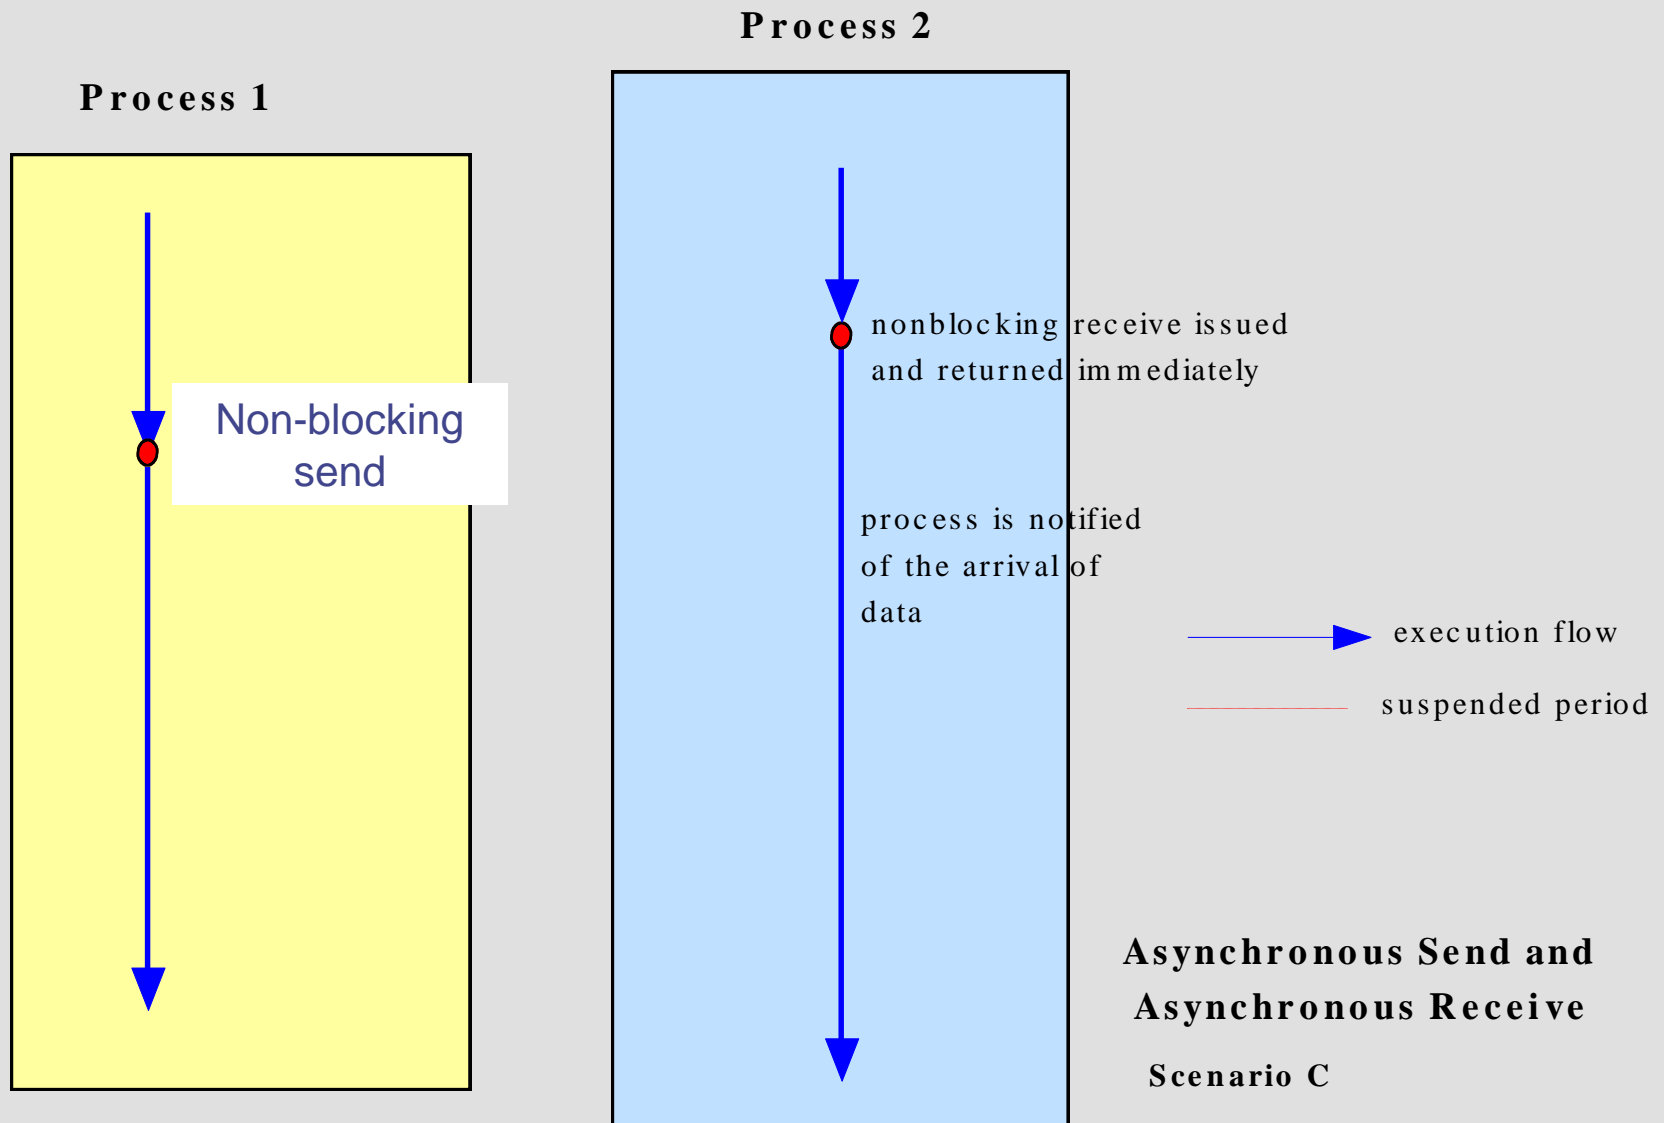
Asynchronous Interaction

- Non-blocking send: sending process continues as soon message is queued.
- Blocking or non-blocking receive:
 - Blocking:
 - Timeout.
 - Threads.
 - Non-blocking: proceeds while waiting for message.
 - Message is queued upon arrival.
 - Process needs to poll or be interrupted.
- Arbitrary processes execution speeds, message transmission delays and clock drift rates
- Some problems impossible to solve (e.g. agreement)

Asynchronous send and synchronous receive

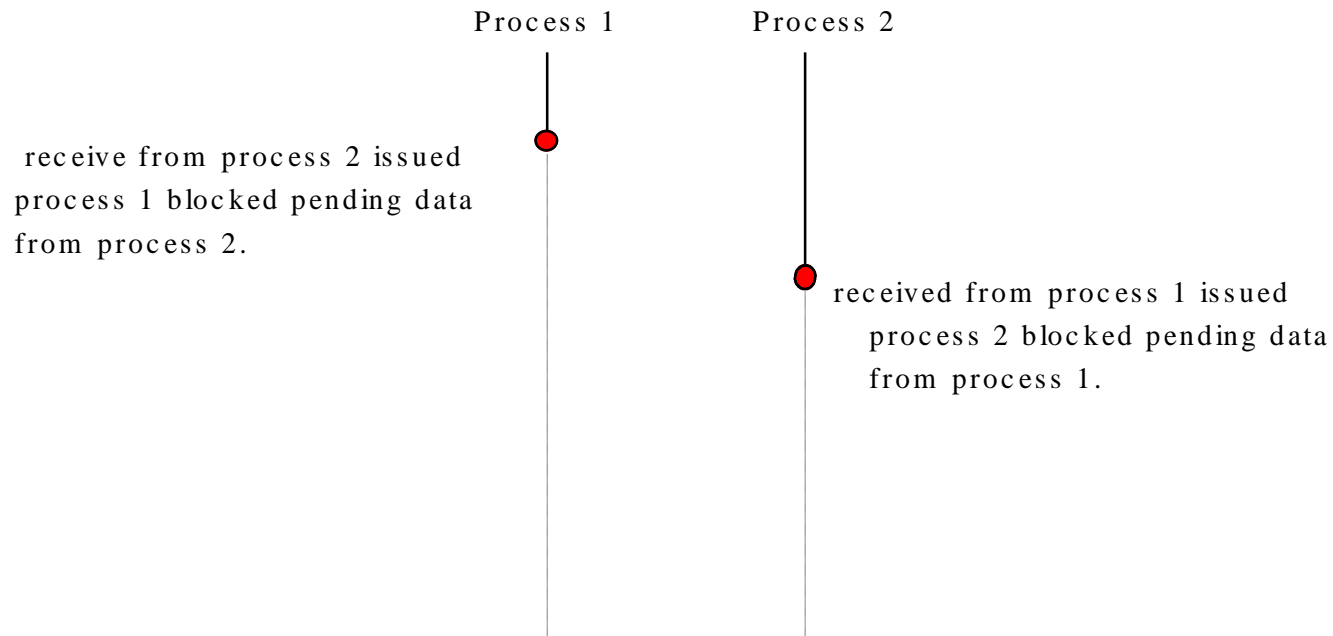


Asynchronous send and Asynchronous receive



Blocking, deadlock, and timeouts

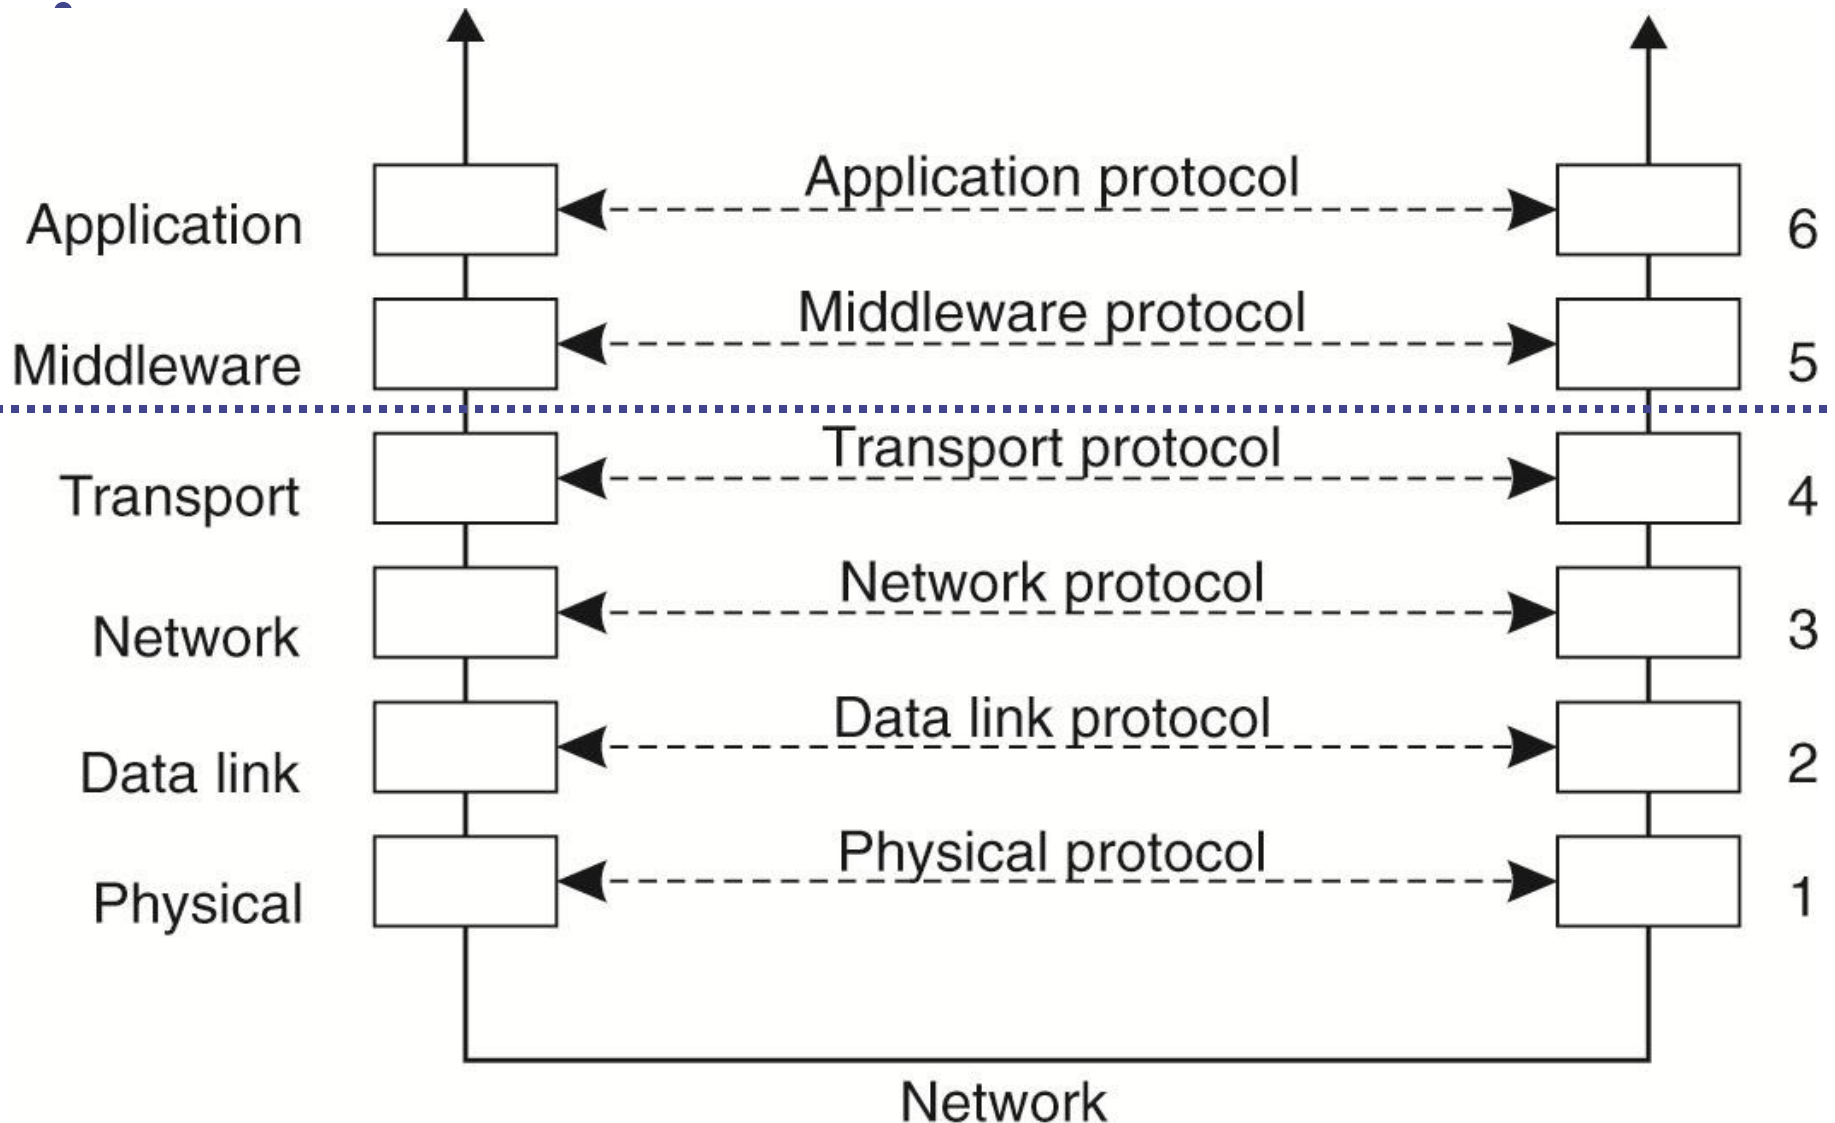
- Blocking operations issued in the wrong sequence can cause deadlocks.
- Deadlocks should be avoided. Alternatively, timeout can be used to detect deadlocks.



Primitive	Meaning
MPI_bsend	Append outgoing message to a local send buffer
MPI_send	Send a message and wait until copied to local or remote buffer
MPI_ssend	Send a message and wait until receipt starts
MPI_sendrecv	Send a message and wait for reply
MPI_isend	Pass reference to outgoing message, and continue
MPI_issend	Pass reference to outgoing message, and wait until receipt starts
MPI_recv	Receive a message; block if there is none
MPI_irecv	Check if there is an incoming message, but do not block

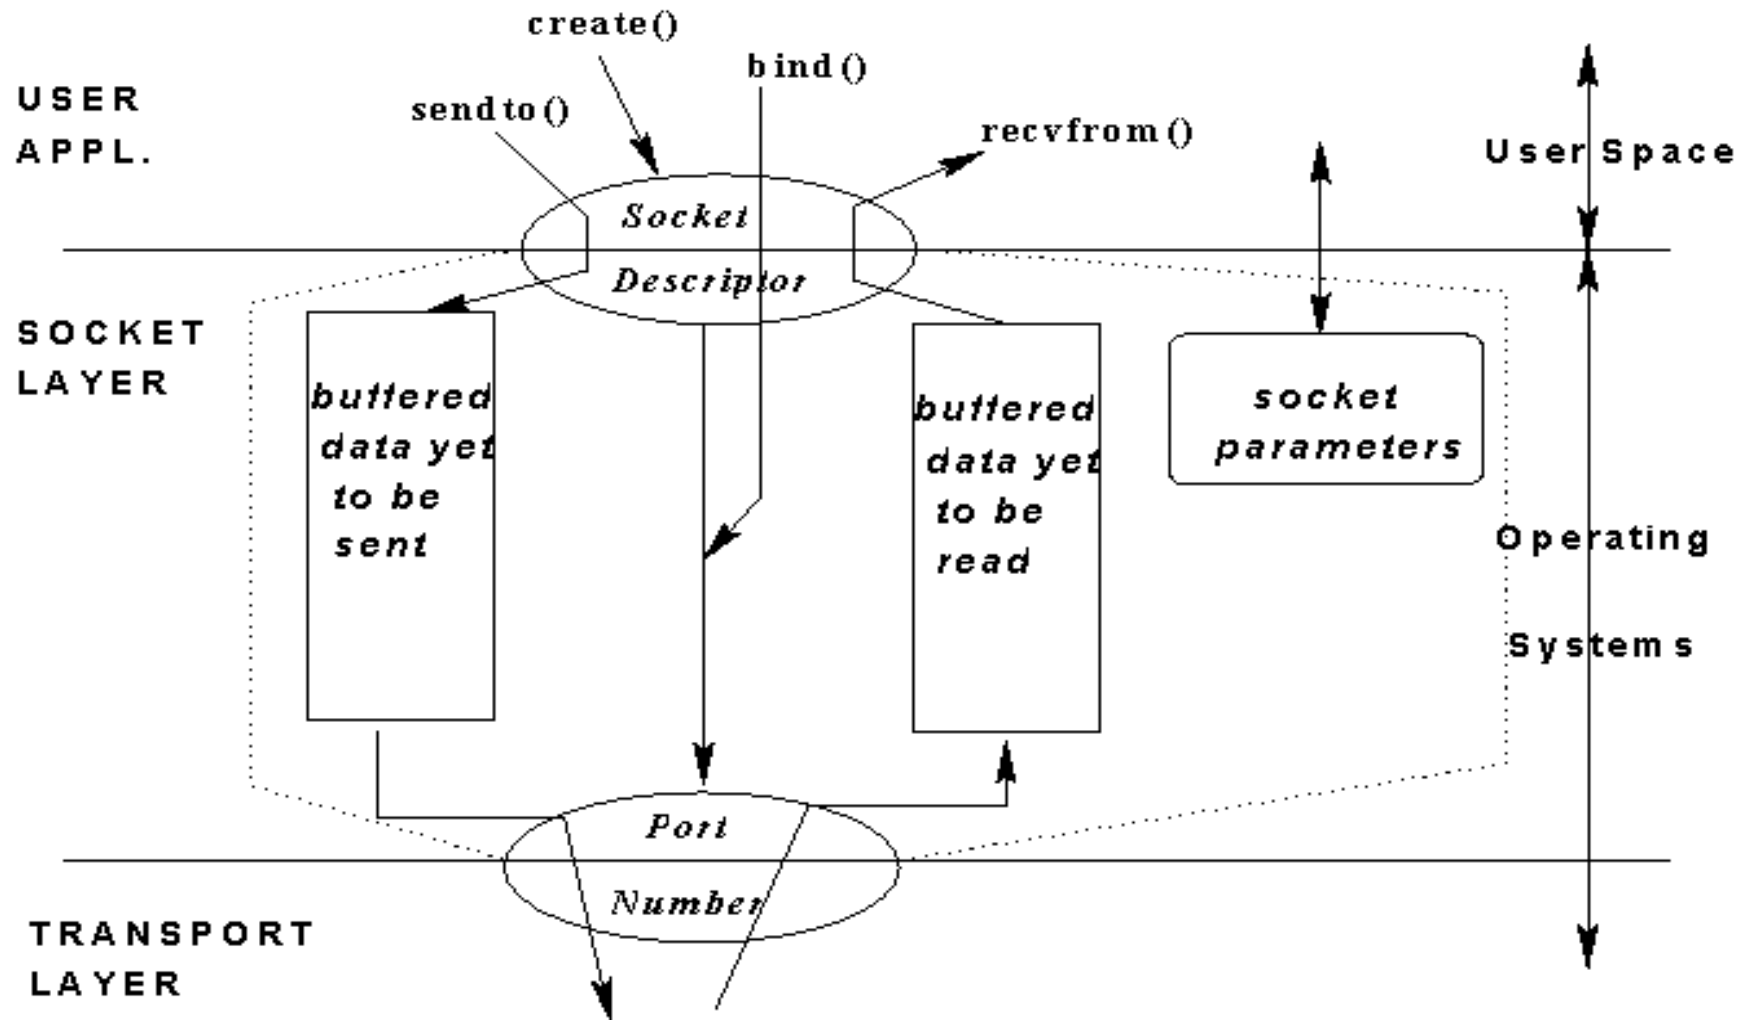
An adapted model communication model

[Tannenbaum]



Sockets: the Transport Layer API

- Sockets provide an API (Application Programming Interface) for programming networks at the transport layer.
- A socket is an endpoint of a two-way communication link between two processes located on the same machine - or located on different machines connected by a network.
- Network communication using Sockets is similar to file I/O
 - Socket handle is treated like file handle.
 - The streams used in file I/O operation are also applicable to socket-based I/O
- Socket-based communication is programming language independent!
- The success of this API is based on its abstraction of all possible used network protocols/underlying network topology.
 - A socket is bound to a port number so that the transport protocol can identify the application that data destined to be sent.



- Recall telephony vs. postal service
 - Service can require a preliminary setup phase, e.g., to determine receiver ! **connection-oriented service**
 - Three phases: connect, data exchange, release connection
 - Alternative: Invocation of a service primitive can happen at any time, with all necessary information provided in the invocation !
connection-less service
 - Note: This distinction does **NOT** depend on circuit or packet switching
 - connection-oriented services can be implemented on top of packet switching (and vice versa, even though a bit awkward)
- Connection-oriented services must provide primitives to handle connection
 - CONNECT - setup a connection to the communication partner
 - LISTEN - wait for incoming connection requests
 - INCOMING_CONN - indicate an incoming connection request
 - ACCEPT - accept a connection
 - DISCONNECT - terminate a connection

Typical examples of services

- ***Datagram service***


- Unit of data are messages (limited length)
 - Correct, but not necessarily complete or in order -
 - Connection-less (*e.g. supported by UDP*)
 - Usually insecure/not dependable, not confirmed
- Application must provide its own reliability“

- ***Reliable byte stream***



- Byte stream
- Correct, complete, in order, confirmed - Processes have some guarantee that messages will be delivered.
 - Possible to build reliability atop unreliable service (E2E argument).
- Sometimes, but not always secure/dependable
- Connection-oriented (*e.g. Supported by TCP*)

TCP supports a “stream of bytes” service

Host A




Host B



...which is emulated using TCP “segments”

Host A




TCP Data


Segment sent when

1. Segment full (MSS bytes),
2. Not full, but times out, or
3. “Pushed” by application.


TCP Data

Host B





- Connection-Oriented Communication



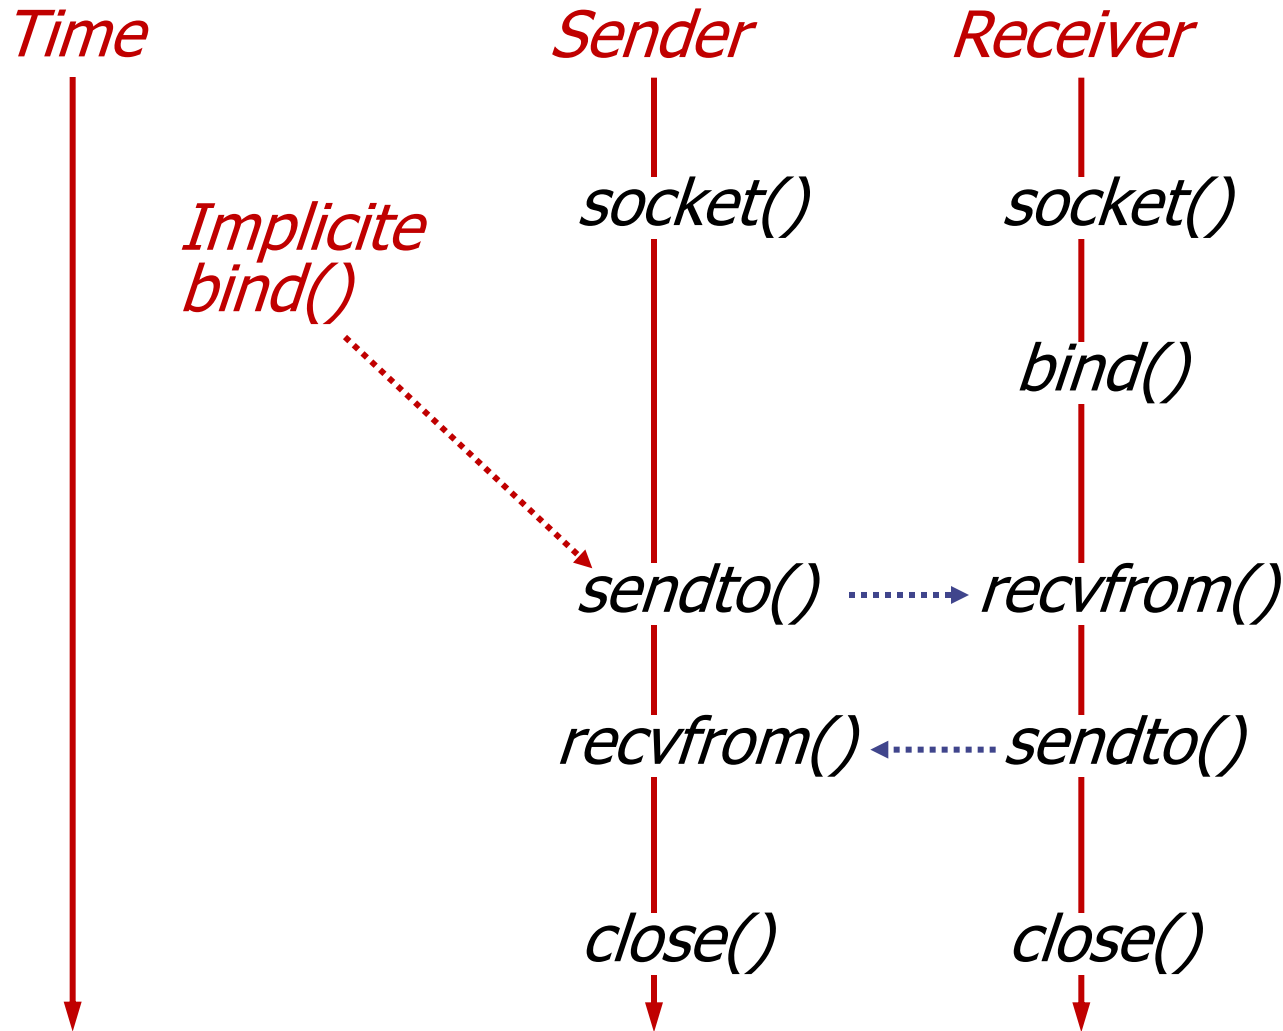
- Connectionless Communication

Berkeley Sockets

Primitive	Meaning
Socket	Create a new communication end point
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection

A VERY GOOD source of information about sockets is the Beej's Guide
<http://beej.us/guide/bgnet/> (legal free download!)

Datagram sockets (see MPGI 4)



- Simplest possible service: unreliable datagrams

Sender

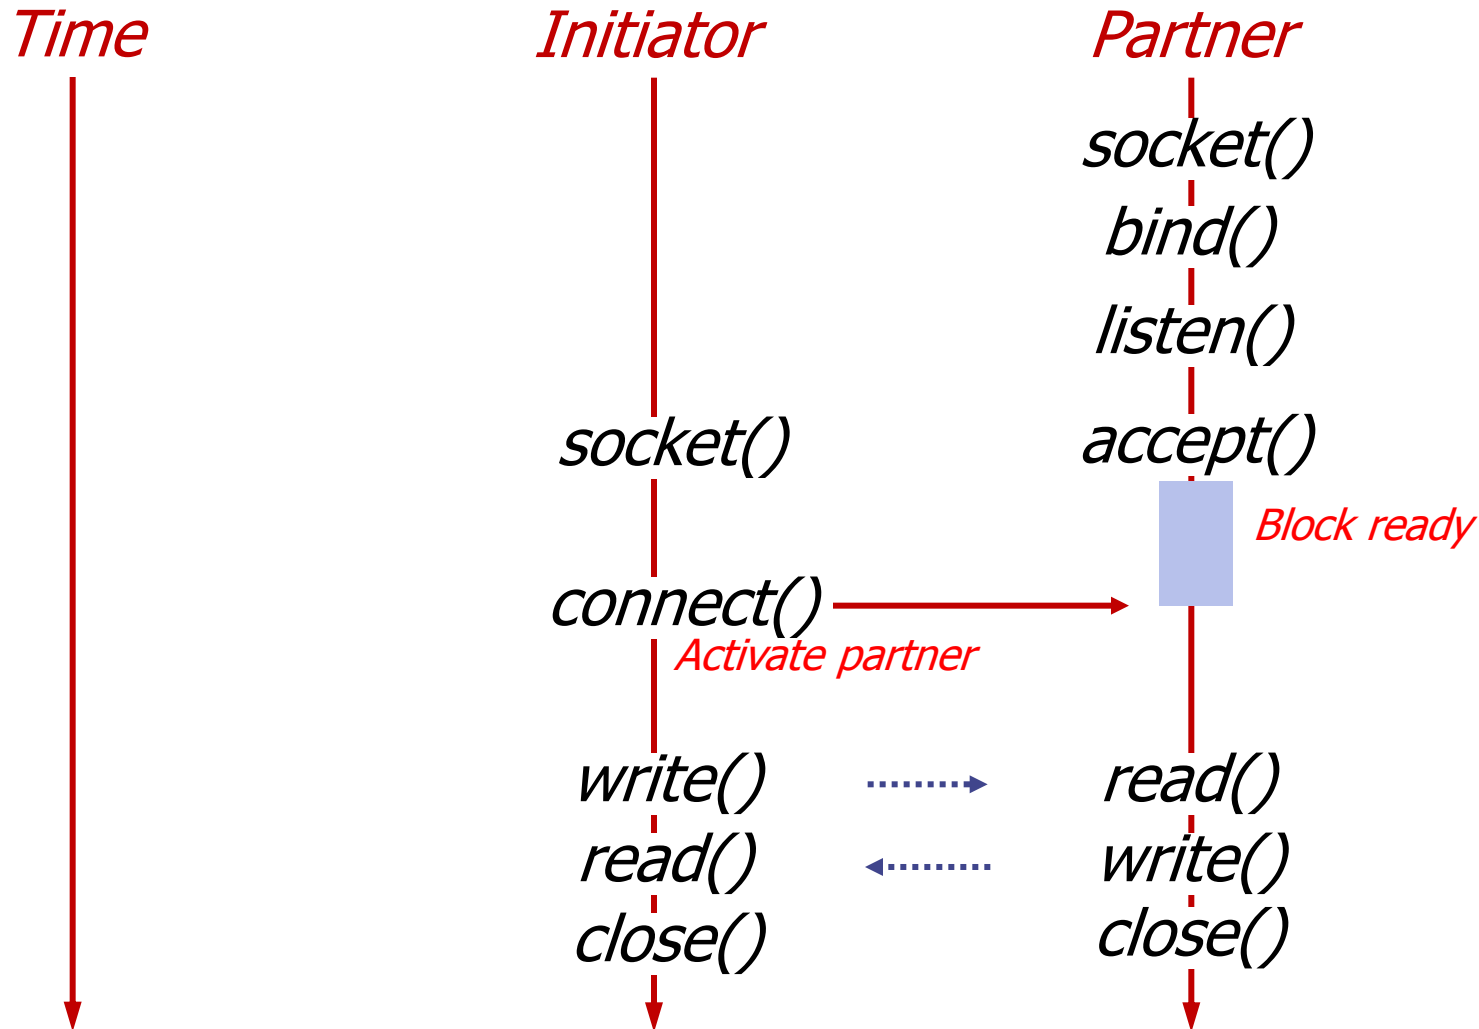
- `int s = socket (...);`
- `sendto (s,
buffer,
datasize,
0,
to_addr,
addr_length);`
- `to_addr` and `addr_length` specify the destination

Receiver

- `int s = socket (...);`
- `bind (s, local_addr,
...);`
- `recv (s,
buffer,
max_buff_length,
0);`
- Will wait until data is available on socket `s` and put the data into `buffer`

- Class *DatagramSocket*
 - *socket constructor* (returns free port if no arg)
 - *send DatagramPacket*, non-blocking
 - *receive DatagramPacket*, blocking
 - *setSoTimeout* (receive blocks for time *T* and throws *InterruptedIOException*)
 - *connect*
 - *close DatagramSocket*
 - throws *SocketException* if port unknown or in use

Stream sockets



- For reliable byte streams, sockets have to be connected first
- Receiver has to accept connection

Initiator (*client*)

- `int s = socket (...);`
- `connect (s,
destination_addr,
addr_length);`
- `send (s,buffer,
datasize, 0);`
- Arbitrary `recv()/send()`
- `close (s);`
- Connected sockets use a `send` without address information

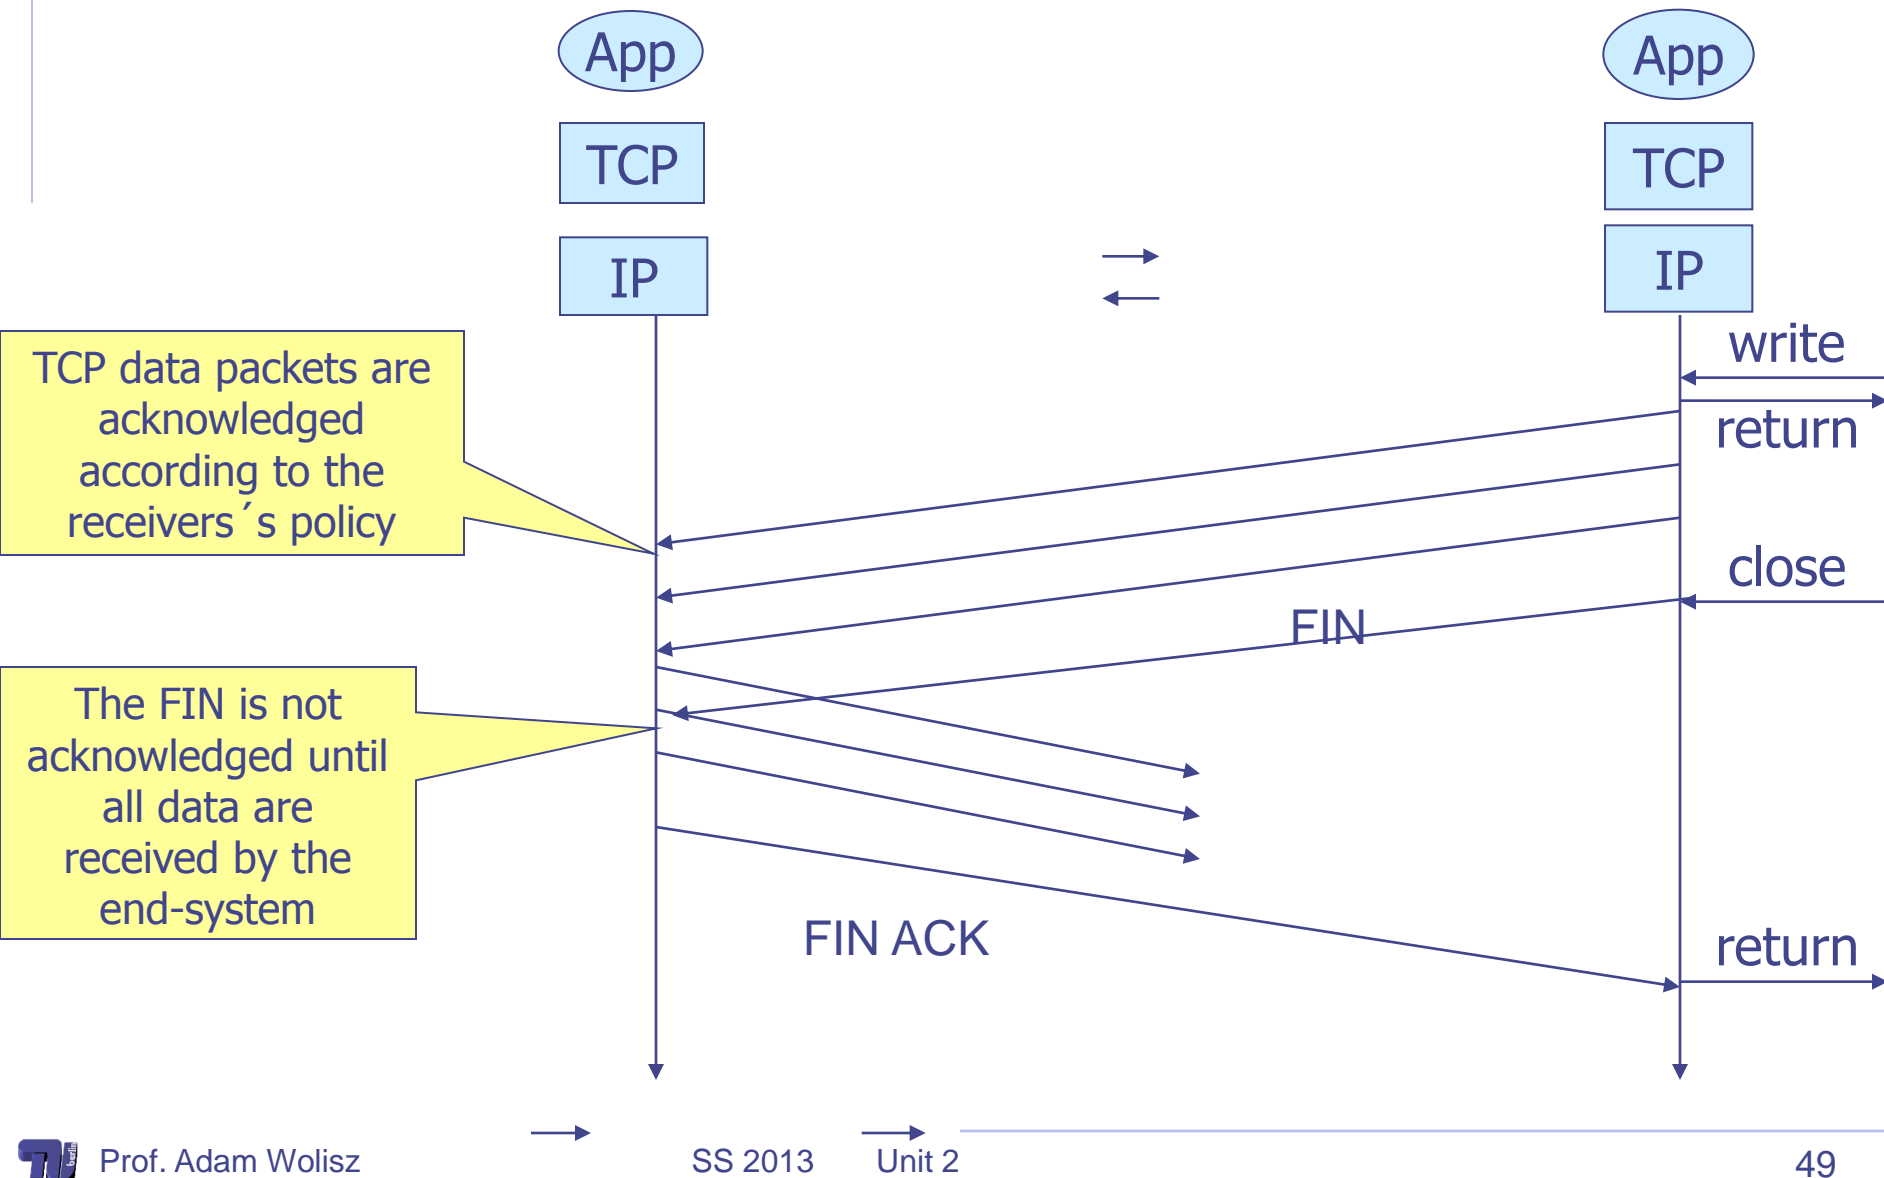
Partner (*server*)

- `int s = socket (...);`
- `bind (s, local_addr, ...);`
- `listen (s, ...);`
- `int newsock = accept (s,
*remote_addr, ...);`
- `recv (newsock, buffer,
max_buff_length, 0);`
- Arbitrary `recv()/send()`
- `close (newsock);`

Java API for Data Stream Communications

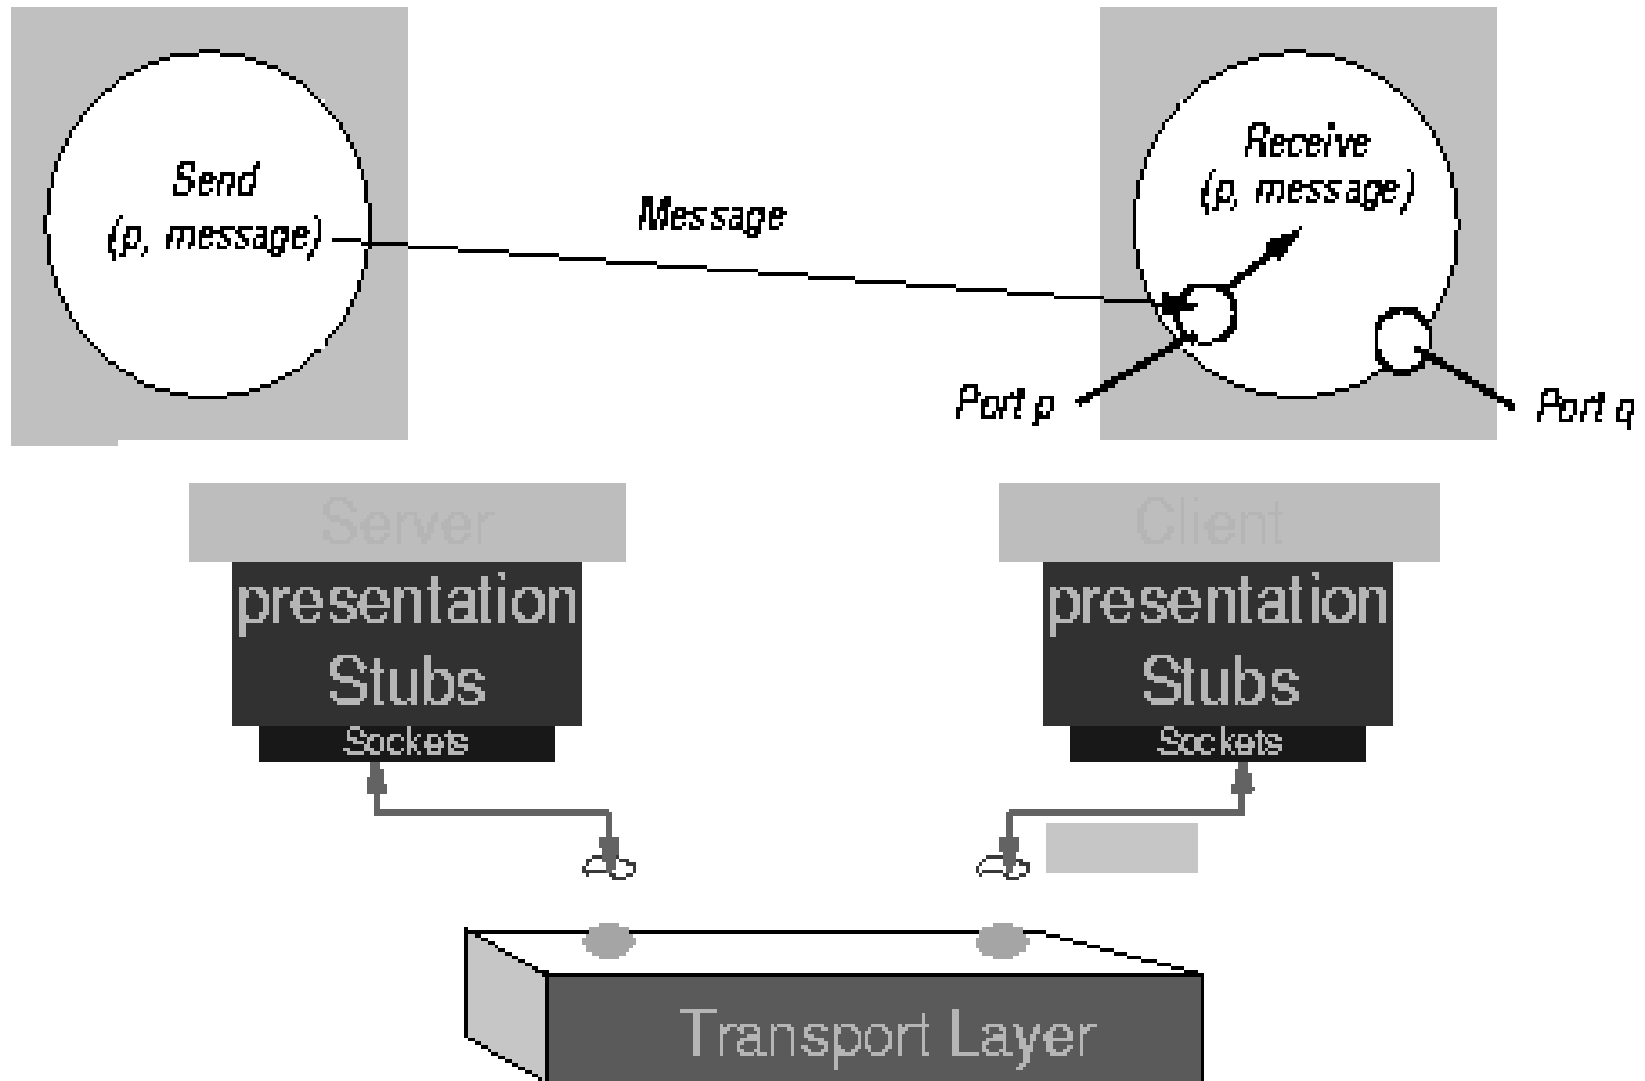
- Class *ServerSocket*
 - *socket constructor* (for listening at a server port)
 - *getInputStream, getOutputStream*
 - *DataInputStream, DataOutputStream*
(automatic marshaling/unmarshaling)
 - *close* to close a socket
 - raises *UnknownHost, IOException*, etc

Reliability of sockets




- For both message and byte stream oriented (set of) service primitives, correctness requirements are important
 - **Completeness:** All data that is sent is eventually delivered
 - **Correctness:** If data is delivered, its is correct, i.e., the data that has been actually sent
 - Messages are not modified, original version is delivered
 - Byte sequence is free of errors
 - **In order:** Byte sequence/sequence of messages is delivered in the order it has been sent
 - **Dependable:** secure, available, ...
 - **Confirmed:** Reception of data is acknowledged to the sender
- Not all requirements are always necessary

Streams of Bits/bytes can be transmitted: so what?
How do we know what is the INFORMATION inside?



Simple example

- Representation of base types
 - floating point: IEEE 754 versus non-standard
 - integer: big-endian versus little-endian (e.g., 34,677,374)




- on a 680x0 CPU, the 32 bit integer number 255 is stored as:
00000000 00000000 00000000 11111111
but an Intel 80x86-CPU stores this as:
11111111 00000000 00000000 00000000

Taxonomy

- Data types

- base types (e.g., ints, floats); must convert
- flat types (e.g., structures, arrays); must pack
- complex types (e.g., pointers); must linearize



- Conversion Strategy

- canonical intermediate form
- receiver-makes-right (an $N \times N$ solution)

Data Conversion

- Two different types of rules are needed:
 - Abstract syntax: a station must define what datatypes are to be transmitted
 - Transfer syntax: it must be defined how these datatypes are transmitted, i.e. which representation has to be used.

Tagged versus untagged data

type = INT	len = 4		value =	417892	
---------------	---------	--	---------	--------	--

Abstract Syntax Notation.1 - ASN.1

- Each transmitted data value belongs to an associated data type.
- For the lower layers of the OSI-RM, only a fixed set of data types is needed (frame formats), for applications with their complex data types ASN.1 provides rules for the definition and usage of data types.
- ASN.1 distinguishes between a data type (as the set of all possible values of this type) and values of this type (e.g. '1' is a value of data type Integer).
- Basic ideas of ASN.1:
 - Every data type has a globally unique name (type identifier)
 - Every data type is stored in a library with its name and a description of its structure (written in ASN.1)
 - A value is transmitted with its type identifier and some additional information (e.g. length of a string).

Definition of Datatypes using ASN.1 (1)

- A data type definition is called „abstract syntax“; it uses a Pascal-like syntax.
- Lexical rules:
 - Lowercase letters and uppercase letters are different
 - A type identifier must start with an uppercase letter
 - Keywords are written in uppercase letters
- ASN.1 offers some predefined simple types:
 - BOOLEAN (Values: True, False)
 - INTEGER (natural numbers without upper bound)
 - ENUMERATED (association between identifier and Integer value)
 - REAL (floating point values without upper or lower bound)
 - BIT STRING (unbounded sequence of bits)
 - OCTET STRING (unbounded sequence of bytes/ octets)
 - NULL (special value denoting absence of a value)
 - OBJECT IDENTIFIER (denoting type names or other ASN.1-objects)

Definition of Datatypes using ASN.1 (2)

- Examples:
 - MonthsPerYear ::= INTEGER
MonthsPerYear ::= INTEGER (1..12)
Answer ::= ENUMERATED (correct(0), wrong(1),noAnswer(3))
- With the following type constructors new types can be built from existing ones:
 - SET: the order of transmission of the elements of a set is not specified. The number of elements is unbounded, their types can differ
 - SET OF: like SET, but all elements are of the same type.
 - SEQUENCE: the elements of a sequence are transmitted in the defined order. They can be of different types. The number of elements is unbounded.
 - SEQUENCE OF: like SEQUENCE, but all elements are of the same type
 - CHOICE: the type of a given value is chosen from a list of types (like a Pascal variant record)
 - ANY: unspecified type

ASN.1 Transfer Syntax (1)

- Some coding rules (the „transfer syntax“) specifies how a value of a given type is transmitted. A value to be transmitted is coded in four parts:
 - identification (type field or tags)
 - length of data field in bytes
 - data field
 - termination flag, if length is unknown.
- The coding of data depends on their type:
 - integer numbers are transmitted in High-Endian Two's complement representation, using the minimal number of bytes: numbers smaller 128 are encoded in one byte, numbers smaller than 32767 are encoded in two bytes, ...
 - Booleans: 0 is false, every value not equal 0 is true.
 - for a sequence type first a type identification of the sequence itself is transmitted, followed by each member of the sequence.
 - Similar rules apply to the transfer of set types

We can now move MEANINGFULL Messages...

but how to use them?

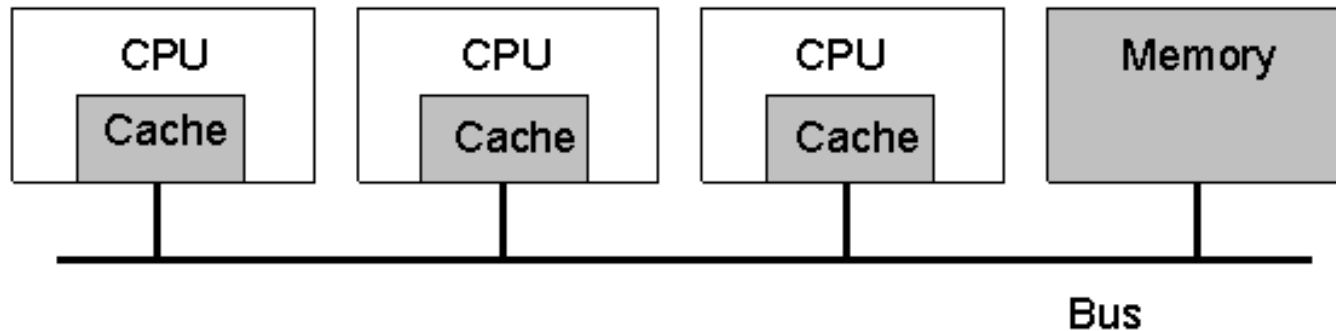
A Distributed System:

- *A distributed computing system consists of multiple **autonomous** processors that do not share primary memory but **cooperate** by sending messages over a communication network.*

Henri Bal/ Colouris

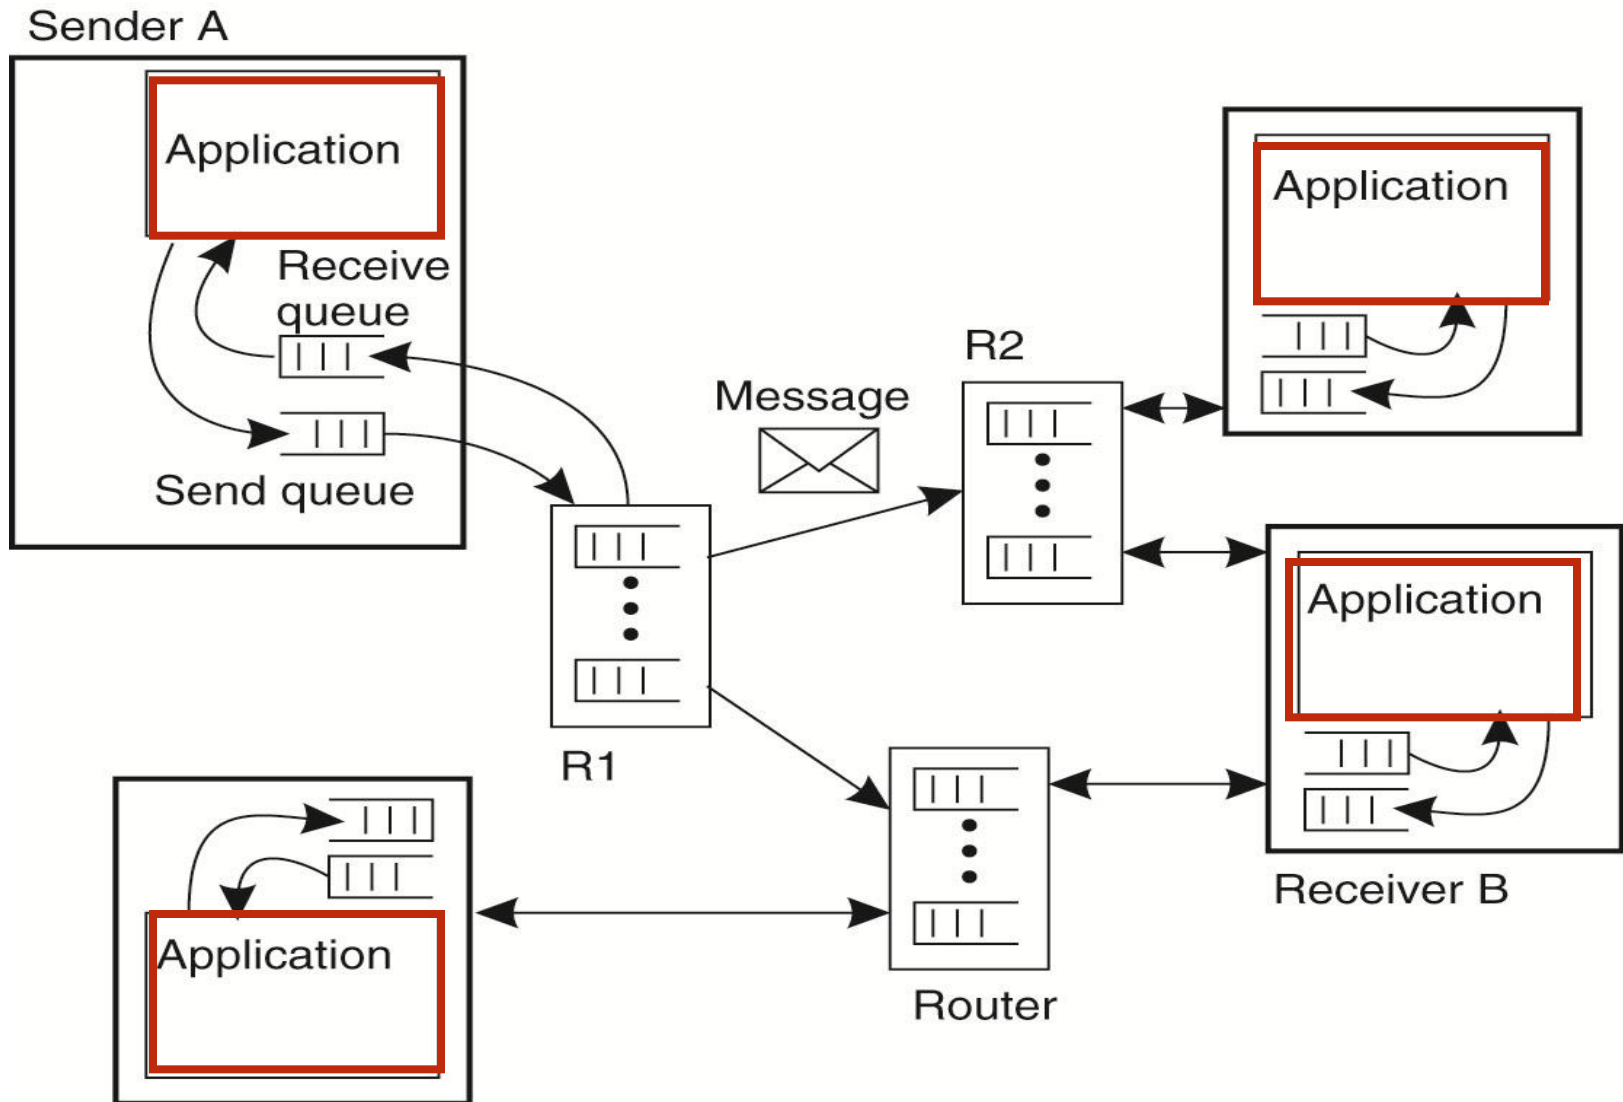
Multiprocessors vs. Distributed systems

- A bus-based multiprocessor.



Typical features of a multiprocessor:

- Physical access to a common memory
- Resources under same management
- Very fast (bus, switching matrix) based local communication



A Distributed System:

- *A distributed computing system consists of multiple **autonomous** processors that do not share primary memory but **cooperate** by sending messages over a communication network.*

Henri Bal/ Colouris

- A distributed system is one in which the **failure** of a computer which you didn't even know existed can render your own computer unusable.

Leslie Lamport

=====

- A distributed system is a collection of **independent** computers that **appears** to its users as a **single coherent** system.

A. S. Tannenbaum

Two important features of Distributed Systems

- **Autonomy**

- A distributed system consists of **autonomous**, independent entities (**usually** cooperative ones, **sometimes** antagonistic ones)
- Each individual entity is – typically – a full-fledged, operational system of its own
- Individual entities might follow local policies, be subject to local constraints...

- **Transparency**


- The fact that the distributed systems is indeed a conglomerate of different (simpler) systems is of no interest and should be of no concern to the user

,

- Consider an action game: player A shoots player B.
 - Each of the players has a computer with a local copy of the game
 - Application of player A creates a bullet entity with a certain heading and velocity. It will then transmit the state of that bullet entity.
 - Upon receiving the state of the bullet, remote applications will start to check whether any entity that they control is hit by the bullet.
- There is network delay between the time A transmits the initial state of the bullet and the time B receives the state.
 - This network delay may be much larger (in the order of 100ms or more) than the amount of time that the bullet needs to hit its target!
 - During this time player B might take actions,
 - shoot at another player C, even though he's already dead.
 - move so that the bullet would not hit him (from his point of view!!!)
- Players A, B, and C may therefore disagree about whether a hit has been scored on B or not!

Dead Reckoning

- Based on ocean navigation techniques
- Predict position based on last known position plus direction
 - Can also only send updates when deviates past a threshold



- When prediction differs, get "warping" or "rubber-banding" effect

The Distributed Stock Exchange...

- Consider a distributed stock exchange...
 - Brokers are distributed all over the world.
 - Each of the brokers runs a computer program
 - Each of the brokers sets his program to:
Buy stocks of Company HOPE if they fall under XXX \$
- The central unit (located at Wall street, NY) defines the value of the HOPE stocks and announces them ...
- Due to different transmission delays Broker John - at Wall Street - will always react quicker than Paul in Frankfurt , or Wendong in Peking, or....
- But the relative advantage of Paul vs Wendong (or vice versa!) might dramatically change depending on traffic conditions....(e.g. local soccer league...)
 - What about Paul causing artificial load on Wendong's computer by sending him numerous superfluous requests from a computer of his girl friend?

Authorizing information exchange

- Smart phones - or Body Area Networks.
- People would like to exchange information locally...
 - With whom? Not with EVERYBODY in your proximity...
- An (older) simple idea:
accelerometers in “wrist watches” for detecting shaking hands... as a criterion for :
 - exchange of contact data between the smart phones?
 - Giving mutual access to the time planner?
 - Some other data?


Handshake data: Red curve- Local; Green curve – Partner.
(Shift of the curves is given by local optical alignment)

- Red square:

The data section providing sufficient match for assuming a joint handshake.

- But: Each of the participant might have a DIFFERENT decision Algorithm...

- ➔ Decisions about data export and data acceptance might be contradictory!!!



Transparency	Description
Access	Hide differences in data representation and how a resource is accessed
Location	Hide where a resource is located
Migration	Hide that a resource may move to another location
Relocation	Hide that a resource may be moved to another location while in use
Replication	Hide that a resource is replicated
Concurrency	Hide that a resource may be shared by several competitive users
Failure	Hide the failure and recovery of a resource

Note: There is a difference between the **FEATURE**, **MECHANISM** and **POLICIES**
Take openness and scaling into account...➔

Different forms of transparency in a distributed system (ISO, 1995)

- Degree to which new **resource-sharing** services can be added & used.
- Requires *publication* of interfaces for access to shared resources.
- Requires *uniform communication* mechanism.
- Conformance of each component to the published standard must be tested and verified.

- A system is said to be scalable if it will remain effective when there is a significant increase in the number of **resources** and **users**:
 - Controlling the cost of resources
 - Controlling the performance loss
 - Preventing software resources running out (e.g., IP addresses)

Characteristics of decentralized algorithms

- No machine has complete information about the system state.
- Machines make decisions based only on local information.
- Failure of one machine does not ruin the algorithm.
- There is no implicit assumption that a global clock (i.e. precise common understanding of time!) exists.

Pitfalls when Developing Distributed Systems

- The network is reliable.
- The network is secure.
- The network is homogeneous.
- The topology does not change.
- **Latency is zero.**
- **Bandwidth (= bit rate!) is infinite.**
- Transport cost is zero.
- There is one administrator.