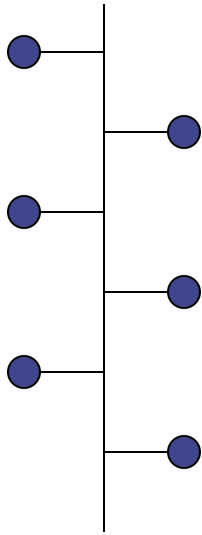


# TechGI IV - ComDiS

## ComDis

Introduction to  
**C**ommunication Networks  
and **D**istributed Systems

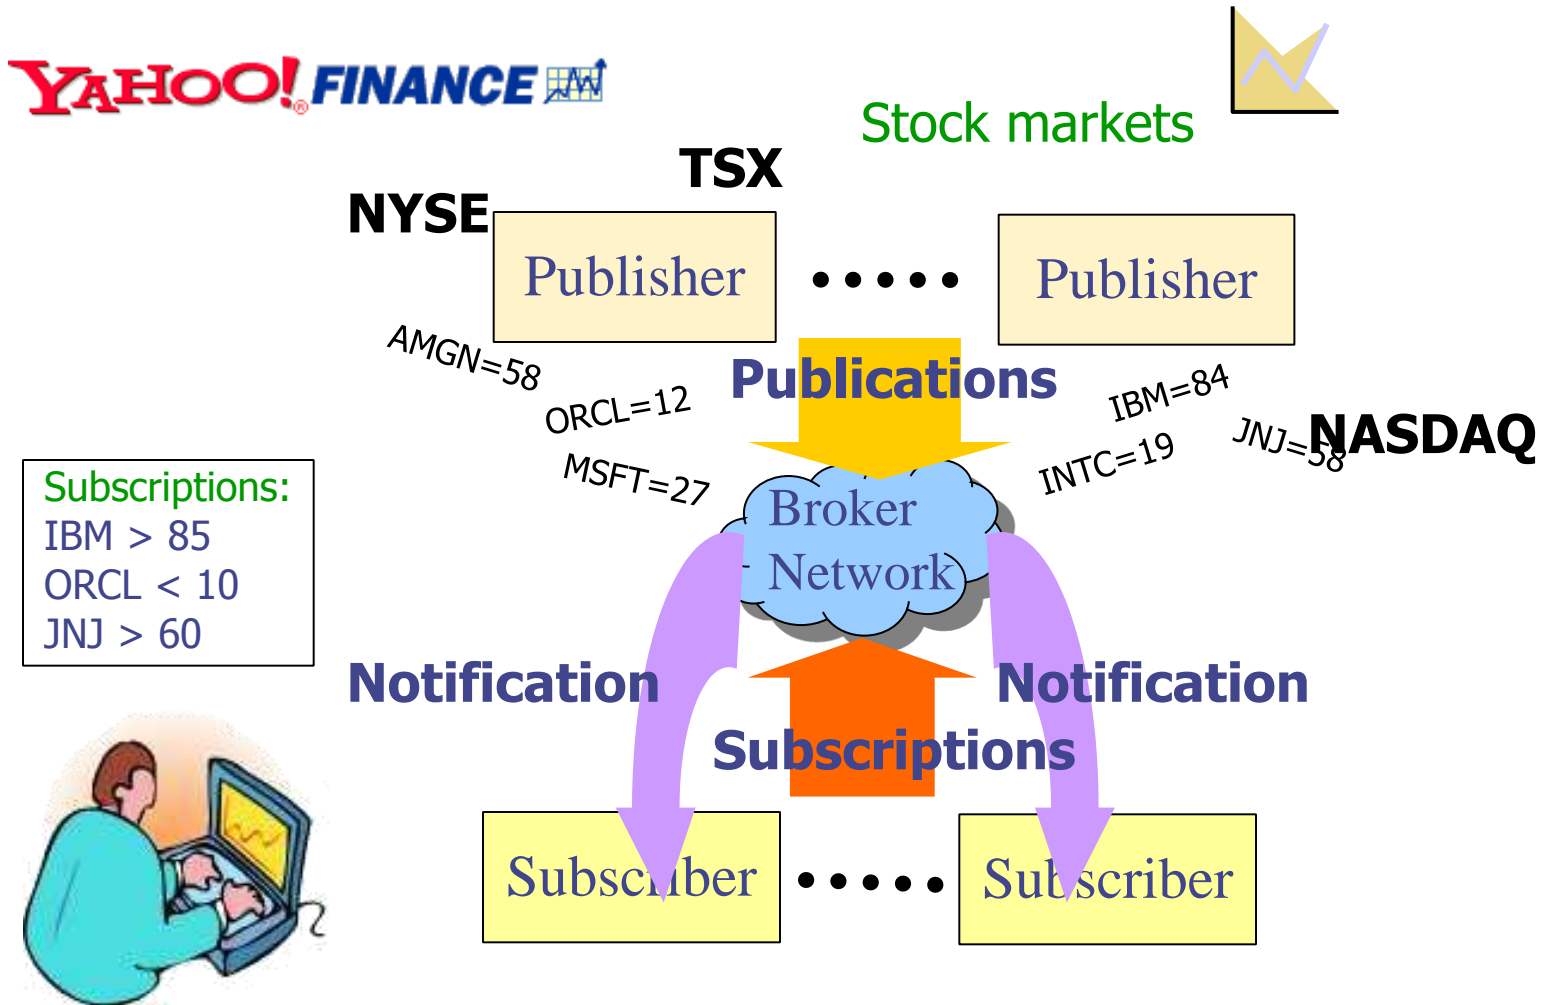


***Complex Communication  
Patterns***

Prof. Dr.-Ing. Adam Wolisz

- **Advanced Communication architectures:**
  - **Publish subscribe**
  - **(Virtual) Shared memory**

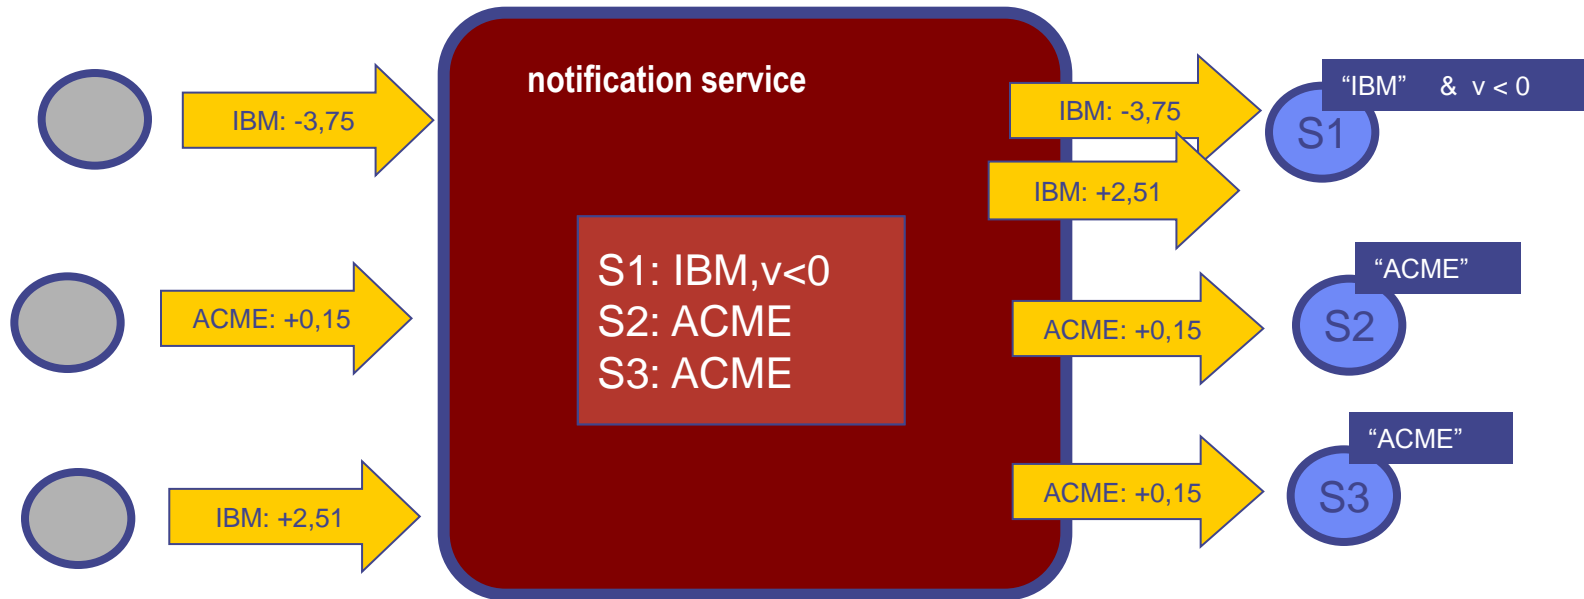
# Publish/Subscribe Model



# Introduction The Publish/Subscribe Paradigm [baldoni]

- Publish/Subscribe (pub/sub): a powerful abstraction for building distributed applications
  - Message-based, anonymous communication
  - Participants are decoupled
    - in space: no need to be connected or even know each other
    - in flow: no need to be synchronized
    - in time: no need to be up at the same time
- Many names for pub/sub systems: notification service, data distribution service
- Applications
  - Stock information delivery
  - Auction system
  - Air traffic control
  - .etc...

# Subscription Models



## **Topic-Based [Oki et al. 93]:**

- events are divided in topics
- subscribers subscribe for a single topic

## **Type-Based [Eugster 2001]:**

- notifications are objects
- type is the discriminating attribute

## **Content-Based [Carzaniga et al 2001]:**


- subscriptions are generic queries SQL-like on the event schema

# Pub/Sub Variants: Topic-based

- Event space is divided in topics, corresponding to logical channels
- Participants subscribe for a topic and publish on a topic
- Receivers for an event are known a priori
- Channel = Group
  - Therefore often exploit network-level multicast
  - Group communication

# Content Based pub-sub: event schema

- Subscriptions and events defined over an n-dimensional *event space* (E.g. `StockName = "ACME"` and `change < -3`)
  - Subscription: conjunction of constraints

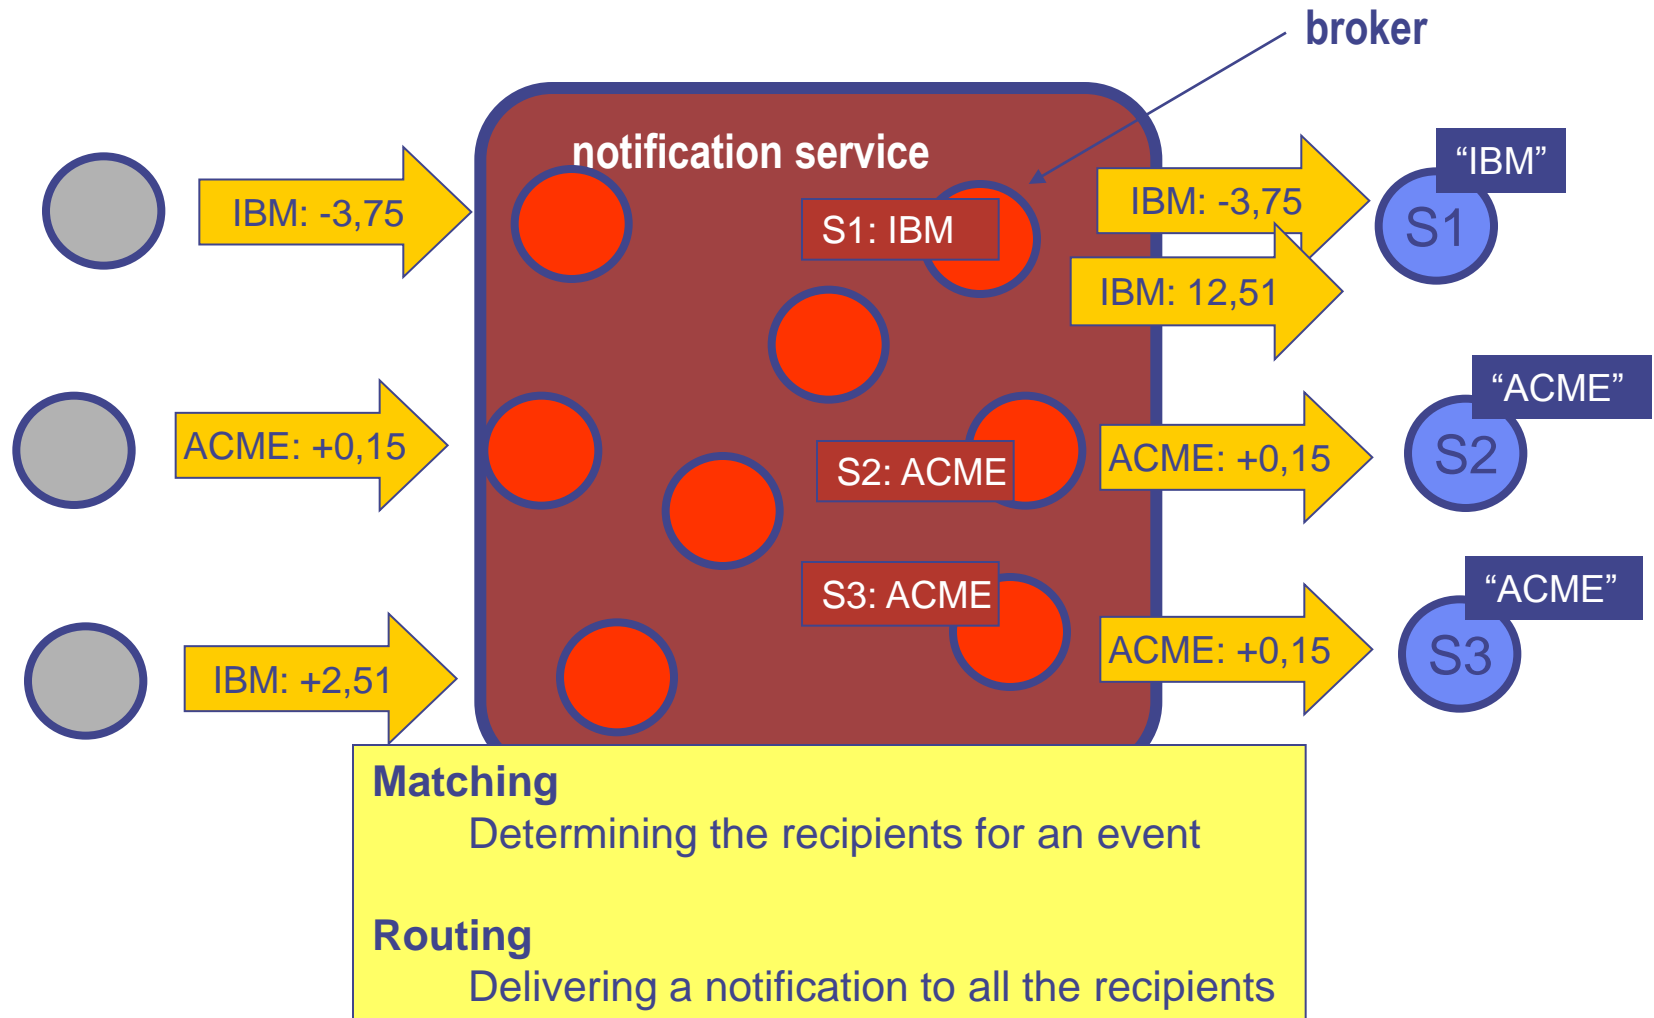


# Topic vs. Content

- Topic-based pub/sub
  - Recipients are known a-priori
  - Many efficient implementations exist: e.g. tib
  - Limited expressiveness
- Content-based pub/sub
  - Cannot determine recipients before publication
  - More flexible
  - More general
  - Much more difficult to implement efficiently




# Distributed Notification Service



# EVENT vs. SUBSCRIPTION ROUTING


- Extreme solutions
- Sol 1 (event flooding)
  - flooding of events in the notification event box
  - each subscription stored only in one place within the notification event box
  - Matching operations equal to the number of brokers
- Sol 2 (subscription flooding)
  - each subscription stored at any place within the notification event box
  - each event matched directly at the broker where the event enters the notification event box

# Intermediate routing solutions




# Intermediate routing solutions

Handling subscriptions: “subscription partitioning” i.e., “where to store  $s$ ”




# Intermediate routing solutions

**Handling subscriptions: “subscription routing” i.e., “how to bring  $s$  there”**




# Intermediate routing solutions

## Handling events




# Intermediate routing solutions

Handling events: “Event Partitioning” i.e., “where to match  $e$ ”



# Intermediate routing solutions

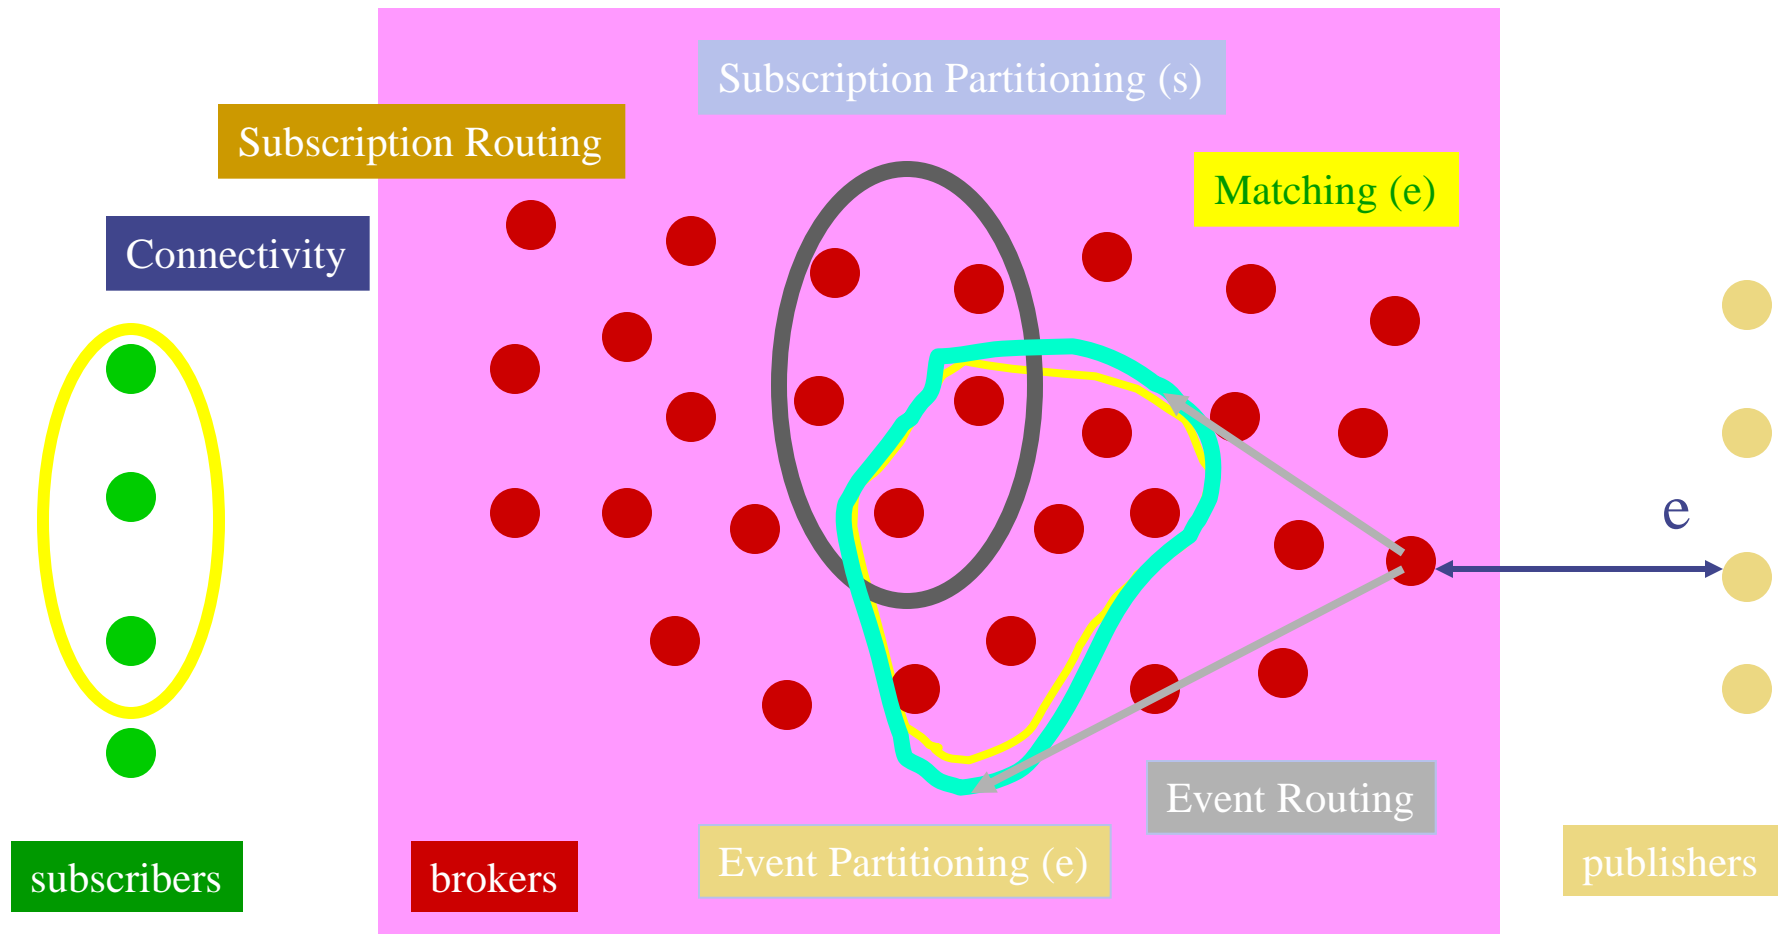
Handling events: “Event Routing” i.e., “How to bring  $e$  there”






# Intermediate routing solutions

Handling events: “Event Matching” i.e., “define the set of recipients of  $e$ ”

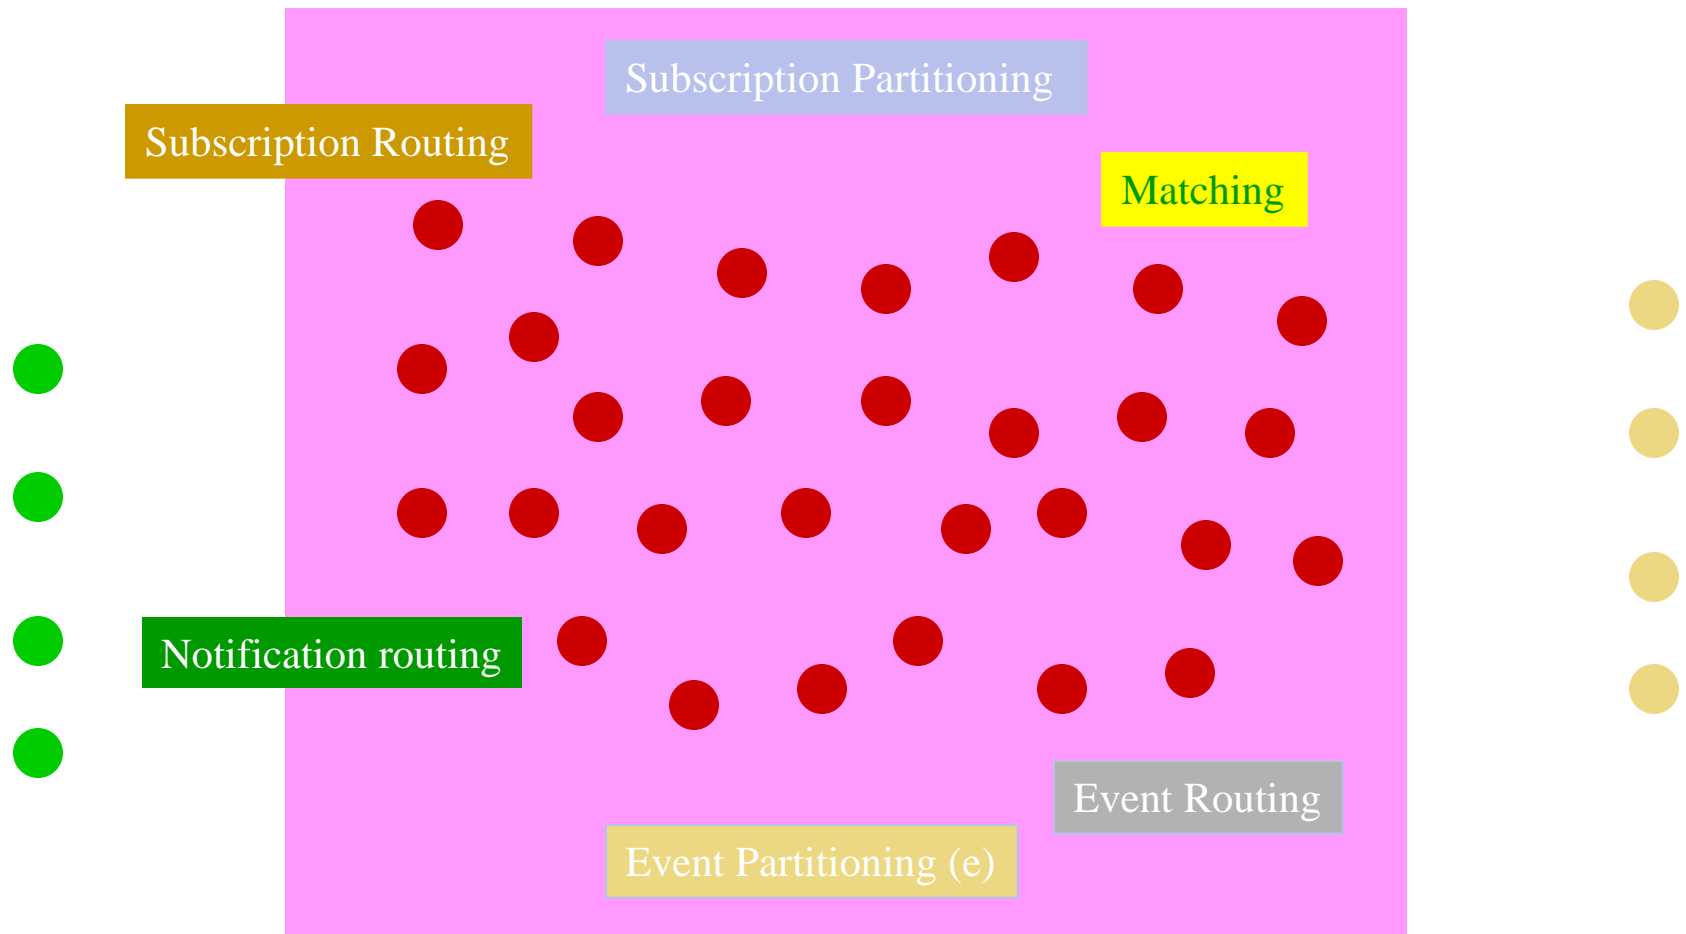


# Intermediate routing solutions

Handling events: “Notification Routing” i.e., “bring  $e$  to the destinations”




# Intermediate routing solutions



# Content-based Routing - example PADRES

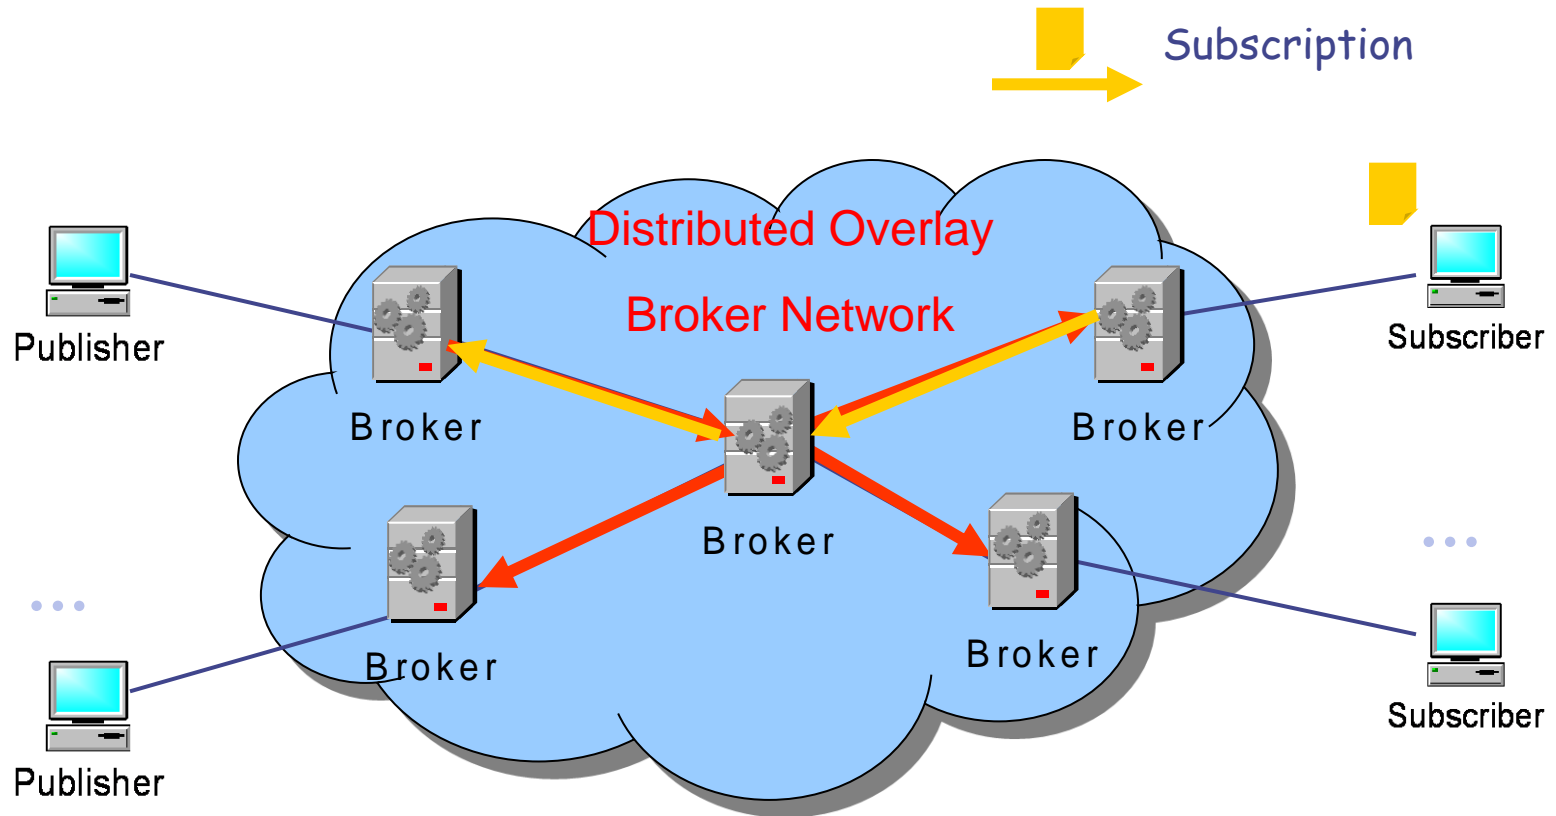
- Advertising



**\*Adopted from SIENA and REBECA**


# Content-based Routing - example: PADRES

- Subscribing



# Content-based Routing - example: PADRES

- Publishing



**\*Adopted from SIENA and REBECA**

- Another abstraction for communication....
- Motivated by development of shared-memory multiprocessors which do share memory.
- Abstraction used for sharing data among processes running on machines that do **not** share memory.
- Processes think they read from and write to a “virtual shared memory”.
- Primitives: read and write.
- OS ensures that all processes see all updates.
  - Happens transparently to processes.

- DSM is an abstraction!
  - Gives programmers the flavor of a centralized memory system, which is a well-known programming environment.
  - No need to worry about communication and synchronization.
- But, it is implemented atop MP.
  - No physically shared memory.
  - OS takes care of required communication.
- For performance, DSM caches data locally.
  - More efficient access (locality).
  - But, must keep caches consistent.
  - Caching of pages for of page-based DSM.
- Issues:
  - Page size.
  - Consistency mechanism.



