

Algorithms for distributed systems

- Overview
 - Distributed mutual exclusion
 - Election algorithms
 - Consensus algorithms

Coordination and matching algorithms

- Coordination and matching algorithms necessary in processes, where e.g.
 - coordination of activities regarding access to jointly used resources such as common objects
 - Agreement on a joint coordinator, for example if the current coordinator fails and need to be replaced ⇒ Implementation of election algorithms in distributed systems
- Motivation
 - Complex devices with several information sources and controlling devices
 - Controlling devices must agree on a certain action, e.g. whether the currently running operation should be continued or aborted
 - Examples
 - Confirm or abort the start procedure of airplane on the runway
 - Activate the brakes in case of danger, e.g. in trains

Coordination and matching algorithms (2)

- Implementation approaches
 - Master node collects all information, evaluates the information and meets the decision
 - Straight forward implementation
 - Single point of failure, which can lead to crash of the entire system
 - Self-organizing process in a set of distributed, involved components
 - Complex implementation
 - Single point of failure avoided
- Examples for characteristic problems
 - Distributed mutual exclusion
 - Election algorithms
 - Consensus algorithms

Distributed mutual exclusion

- Critical sections: Chain of operations, where a concurrent processing may cause errors
- Examples
 - Object update (e.g. write access to a file): Lock the file during the first access, further write operation after write lock release by process currently updating the object
 - Peer-to-Peer coordination while using joint resources
 - Mobile ad-hoc networks \Rightarrow solely one node is allowed to transmit on a jointly used channel
 - Real life: Number of free parking slots in a parking house with several entrances and exits
 - One process per entrance or exit
 - Common counter for the entire house
 - Consistent counter update by distributed processes necessary
 - We need an algorithm for distributed mutual exclusion without application of a master node

Mutual exclusion with semaphores

- Problem well known in system software
- Application of binary semaphores to implement mutual exclusion
 - Two states (locked, free) to protect critical sections
 - Critical section free, then access granted, otherwise the process is blocked (e.g. `sleep()`) and inserted into the queue assigned to the semaphore
- Solution not applicable to distributed systems, as no shared memory available \Rightarrow approach with messages needed

Mutex=Mutual Exclusion

Semaphore *Mutex*=1

Process *P1*

Process *P2*



P(*Mutex*)

P(*Mutex*)



Critical
Section


V(*Mutex*)

V(*Mutex*)



Mutual exclusion with master and message passing

- Simplest solution: Using master node
 - Master as a superior unit grants or rejects access into the critical section
 - Process requests access by sending a request message to master node
 - Master queues all requests according to certain criteria (FCFS, priority, ...) and sends reply message to the first element
 - Reply message contains a token as a ticket for the critical section
 - After leaving critical section, process sends the token to the master node and commits the critical section is free for the next process



Implementing mutual exclusion with master node

- Needed functions
 - Inter-process communication with `send()` and `receive()`
 - `send(destination, &message)`
 - `receive(source, &message)`
 - Indirect addressing using a data structure called mailbox = waiting queue for temporary message storage
 - Implementation with standard OS system calls
- Coordination procedure
 - Master queues all requests and sorts according to given criterion
 - First process receives reply message with the token
 - All other requests are blocked until the token is released
 - Then the next process in the queue receives the token

Implementing mutual exclusion with master node(2)

All processes use the mailbox *Mutex* for sending and receiving messages

```
const n=... number processes
main()
{
    createmailbox(mutex);
    send(mutex,&init);
    // Initial state with one token
    // Independent processes
        P(1);
        P(2);
        ...
        P(n);
    // in the distributed system
}
```


```
void P(int i) {
    message msg;
    while(TRUE) {
        receive(mutex, &msg);
        // Token received as
        // necessary prerequisite
        // to access critical section
        <critical section>;
        // Mailbox empty
        send(mutex,&msg);
        // Return the token to mailbox
        // ⇒ critical section free again
    }
}
```


Evaluation criteria

- Performance of an algorithm for mutual exclusion is evaluated according to following criteria
 - Used bandwidth: Proportional to number of sent messages
 - How long a client is deferred while waiting for entrance of exit from the critical section
 - System throughput
 - Efficiency regarding all processes requesting access to the critical section
 - Synchronization delay
- Example: Performance of the algorithm with a master node
 - Two messages needed to enter the critical section, even in case of a free critical section \Rightarrow Delay equals the round trip time for messages
 - Server as bottle neck

Ring based algorithm

- Mutual exclusion between N processes without additional master process
 - Processes arranged in a logical ring structure
 - Process P_i with communication channel to the next process $P_{(i+1) \bmod N}$
- Approach
 - Entry token circulates in the ring
 - Processes without entry request pass the token to their neighbor
 - Process requesting entry keep the token and enter the critical section. After leaving the critical section, the token is passed to the next process in the ring
- Evaluation
 - High bandwidth consumption
 - Entry delay: $0 \dots N$ messages
 - Synchronisation delay: $1 \dots N$ messages



Multicast algorithm with logical time stamp

- Considering peer processes
 - Process requesting entry sends multicast message (to all other peers)
 - Access granted, if and only if all other peer processes reply to the request
- Prerequisites
 - Each process P_i has a counter with logical time
 - Messages include tuples $\langle T, P_i \rangle$ with T as time stamp and P_i as sender ID
 - Each process holds the current status in the variable state
 - RELEASE: outside the critical section, no entry requested
 - WANTED: entry requested
 - HELD: process already entered the critical section
- Procedure
 - On request message, all processes test their status: if all processes in state RELEASE, the processes reply immediately to the request \Rightarrow access granted
 - At least one process in state HELD \Rightarrow no reply \Rightarrow access temporarily rejected

Multicast algorithm with logical time stamp (2)

- (Nearly) simultaneous requests
 - Two or more processes
 - request access \Rightarrow process with smaller time stamp receives all $N-1$ replies first \Rightarrow Access granted
 - Identical Lamport time
 - total ordering using PIDs
- Pro
 - Synchronisation delay of transfer time for a single message
- Con
 - Expensive algorithm regarding bandwidth

On initialization

state := RELEASED;

To enter the section

state := WANTED;

Multicast request to all processes;

// request processing deferred here

T := request's timestamp;

Wait until (nr of replies received = (N - 1));

state := HELD;

On receipt of a request $\langle T_i, p_i \rangle$ at p_j ($i \neq j$)

if (state = HELD or

(state = WANTED and $(T, p_j) < (T_i, p_i)$))

then

queue request from p_i without replying;

else

reply immediately to p_i ;

end if

To exit the critical section

state := RELEASED;

reply to any queued requests;

Leader election algorithms


- Goal: elect a leader that will take over a certain function
 - Elect a (replacement) master node
 - Important: all nodes accept the decision
- Basic assumptions
 - Each participant has a unique identifier
 - Goal is to choose that member with the largest identifier as leader
 - Set of all identifiers unknown to all participants
- Fault assumptions
 - Processes may or may not fail, may behave in a hostile fashion
 - Messages may or may not be lost, corrupted, ...
 - Different algorithms can handle different fault assumptions
- Time assumptions
 - Synchronous time model – all processes operate in lock-step, bounded message transit time?
 - Asynchronous model – no such bounds available?

Leader election algorithms – Concurrency

- Assumptions
 - Any process is allowed to start an election \Rightarrow but one process is allowed to start only one election process at time
 - N processes can start N concurrent elections \Rightarrow the final decision must be consistent despite the concurrency
 - Process P_i is at any time either a participant of the election or the process is not involved in the election
 - Each process receives a unique identifier ID \Rightarrow process with the largest ID wins the election
 - Each process has a variable elected , where the ID of the elected process is stored \Rightarrow Special symbol denotes an undefined value, e.g. election not finished yet

Ring-based election algorithms

- Available for processes arranged in a logical ring
 - Each process has a channel to his direct neighbors
 - Messages are sent in one direction, asynchronously
- Goal: Elect the processes with the largest ID as coordinator
- Initial state
 - Process 17 starts the election, changes the variable state to participant, sends an election messages with his own ID
 - Each receiving process compares the ID contained in the message with his own ID



Ring-based election algorithm (2)

- Result of the comparison
 - Received ID larger than the own ID \Rightarrow change status to participant, pass the message unchanged to the next neighbor
 - Received ID smaller than the own ID and status "no participant" \Rightarrow change status to participant, insert own ID in the message, pass the message to the next neighbor
 - Received ID identical with own ID (and already participant) \Rightarrow current process has the largest ID \Rightarrow Election finished \Rightarrow Coordinator set the variable to no participant and gives his ID in an elected message to the next neighbor
 - Receiving processes change status to "no participant", note the coordinator and pass the message to the next neighbor
- A process has to receive its own ID before sending elected messages
- Important: unique IDs, so two elected processes also in case of concurrency not possible
- Worst-case performance: $3N-1$ messages
- Communication failures or node crashes stop the entire algorithm \Rightarrow little use in real world applications

Bully algorithm


- Elect a process from P_1, P_2, \dots, P_N with the largest ID as coordinator
- Prerequisites
 - Reliable communication, individual processes may crash
 - Synchronous communication \Rightarrow Node failure detection with time outs
 - Each process knows all processes with larger ID than its own and is able to communicate with those
- Used message types
 - Message election: call for election
 - Reply message answer: reaction to election call
 - Coordinator message: call of the elected coordinator
- Setting time outs
 - Estimate the maximal transmission delay T_{trans} and the maximal processing delay $T_{process}$
 - Upper bound $T = 2T_{trans} + T_{process}$
 - Time out expired \Rightarrow notify node failed

Bully algorithm: procedure

- Process with largest ID elects itself to coordinator
- Process P_i with non-maximal ID calls for election
 - Election message to all processes with higher IDs
 - Wait for answer during time T
 - No message received \Rightarrow the process that called the election declares itself to coordinator and sends the coordinator message
 - Receives a process P_j with $ID_j > ID_i$ a message, then the process sends a reply message to P_i . P_i has to wait for an additional interval T' for the coordinator message. If no message received during this time span, a new election is started
 - Process P_j starts a new election and repeats the workflow until the process with the largest ID was determined
- If the former, temporarily failed coordinator is restarted, it starts a new election. If the process still has the largest ID, then this process will be elected as coordinator again

Bully algorithm: Example

- Process P1 detects coordinator (P4) failure and starts an election (Phase 1)
- P2 and P3 send a reply answer to P1 and start their own elections
- P2 receives reply from P3, P3 receives no replies \Rightarrow P3 has the largest ID and declares itself to coordinator (Phase 2)
- Assume: Before P3 is able to send the coordinator message, P3 fails
- P1 waits until interval T' expired without receiving a coordinator message \Rightarrow New election is started
- P2 elected to coordinator




Consensus algorithms

- Consensus decision-making = group decision making process that not only seeks the agreement of most participants, but also to resolve or mitigate the objections of the minority to achieve the most agreeable decision
 - Redundant devices agree on GO or NO GO
 - Bank transfer: involved nodes agree to commit the same amount on both accounts (negative and positive)
 - Entering a critical section \Rightarrow Decision on elected process
- Number of specialised protocols for sub-problems exists
- Looking at the general problem and solution approaches
- Possible failure sources
 - Errors during communication \Rightarrow Loss or manipulation of messages
 - Failed processes lead to unpredictable behavior
 - Fail-stop failure: Process stops \Rightarrow Failure discovery with time-outs
 - Byzantine failure: process delivers wrong results \Rightarrow incorrect messages are produced \Rightarrow Threat for the integrity of the entire system

Definition

- Each process P_i is in an undefined state and votes for a decision v_i , which is part of the set with all possible decisions D ($i=1,2,\dots,N$)
 - Processes exchange messages with their votes
 - Decision based on the selected algorithm
 - Each process stores the agreed value in the decision variable d_i
 - Change of condition in dedicated state, d_i must not be changed any more
- Example
 - 2 processes vote for proceed, 1 process votes for abort
 - Last process crashes \Rightarrow decision for proceed



Consensus for three algorithms

Demands


- For correct execution of a consensus algorithms, following must be fulfilled
 - Termination: Each correct process sets the decision variable at some point of time
 - Agreement: decision value is identical for all correct processes: if P_i and P_j are in dedicated state, then $d_i = d_j$ ($i, j = 1, 2, \dots, N$)
 - Integrity: If all correct processes voted for the same value, then each process adopted this value in the dedicated state
- Problems?
 - Consensus in fail-safe systems: wait, until all votes arrived – including the own vote – and analyse the majority \Rightarrow In worst case, no majority can be achieved and the state remains undefined \Rightarrow Additional decision measures necessary
 - Failures possible \Rightarrow not sure that all votes arrive or arrive correct
 - Reaction and termination of the consensus algorithm must be guaranteed

Example: Byzantine Generals' Problem

- Agreement problem
 - Generals of the Byzantine Empire's army must decide unanimously whether to attack some enemy army
 - Problem is complicated by
 - Geographic separation of the generals, who must communicate by sending messengers to each other
 - Presence of traitors amongst the generals.
 - Traitors can act arbitrarily in order to achieve the following aims
 - Trick some generals into attacking \Rightarrow Force a decision that is not consistent with the generals' desires, e.g. forcing an attack when no general wished to attack \Rightarrow failed processes
 - Catch a messenger \Rightarrow some generals are not able to make their mind \Rightarrow failed communication channels
 - If the traitors succeed, any resulting attack is doomed, as only a concerted effort can result in victory

Byzantine Generals in a synchronous system

- Simple variant: order by commander transmitted to other generals
- Prerequisites
 - Up to f of N processes may fail or produce false results \Rightarrow failed process can send a message with any value (also wrong values) any time
 - Missing messages recognized via time-outs
 - Private communication channels \Rightarrow other processes can't read the transmitted data
- Impossible mission in case of three available processes \Rightarrow if one of the processes fails, no solution possible (Extension for $N \leq 3f$)




- Scenario

- Prefix = Message source
 - $:$ = „says“, $3:1:u = 3$ says 1 says u
 - Commander gives the same order to both sub-commanders, P_3 transmits a wrong order
- \Rightarrow P_2 is not able to detect whether the order by the commander or by the sub-commander P_3 is wrong


Byzantine Generals in a synchronous system (2)

- Scenario: Commander failed
 - Different values to sub-commanders
 - P2 receives different values
 - ⇒ Same situation as in case of failed P3




- Solution for $N = 4$ and 1 failed process
 - Correctly working commanders have an agreement in two message passing intervals
 - Commander sends the value to each sub-commander
 - Each sub-commander sends the received value to the peered sub-commanders
 - Each sub-commander receives the value from the commander and $N-2$ values from peered sub-commanders
 - Commander failed \Rightarrow all sub-commanders may have the same value set, no failure is recognized, or they may have different values \Rightarrow failure detected
 - One of the sub-commanders failed \Rightarrow Detection by different value sets

Example for 4 sub-commanders



- Failed sub - commander
 - \triangleright P2 determines $\text{majority}(v,u,v)=v$
 - \triangleright P4 determines $\text{majority}(v,v,w)=v$
- Both processes detect failed sub-commander



- Failed commander
 - \bullet P2 determines $\text{majority}(u,v,w) = \perp$
 - \bullet P3 determines $\text{majority}(u,v,w) = \perp$
 - \bullet P4 determines $\text{majority}(u,v,w) = \perp$
- \perp = special symbol, no majority possible

Byzantine Generals in asynchronous system

- There is no algorithm that guarantees consensus in asynchronous system, if processes may fail
 - Processes can answer to messages at any point of time \Rightarrow differentiation between slow/delayed and failed processes not possible
 - Proof for non-existence of such an algorithm by Fischer et al.
- Approaches for work-around
 - Using partially synchronous algorithms: relaxed synchronous distributed systems, but still with well-defined upper and lower bounds
 - Masking faults: process data is stored persistently and restored in case of process failure \Rightarrow hiding problems by controlled multiple execution, e.g. in case of transactions
 - Failure detectors: Processes agree, if some of them did not answer for a certain time and send no keep alive signal, then they are declared failed and are removed from the future decision process